SPIR-V Extended Instructions for GLSL

John Kessenich, Google

Version 1.00, Revision 5

December 26, 2017

Contents

1 Introduction

This specifies the GLSL.std.450 extended instruction set. It provides instructions for the GLSL built-in functions that do not directly map to native SPIR-V instructions.

Import this extended instruction set using an **OpExtInstImport** "GLSL.std.450" instruction.

2 Binary Form

Documentation form for each extended instruction:

Extended Instruction Name Instruction description. *Result Type* will describe the *Result Type* for the **OpExtInst** instruction. *Number* is the extended instruction number to use in the **OpExtInst** instruction. *Operand 1*, *Operand 2*,... are the operands listed for the **OpExtInst** instruction. Any **Capability** restrictions.			
Number	*Operand 1*	*Operand 2*	...

Extended instructions:

Round Result is the value equal to the nearest whole number to *x*. The fraction 0.5 will round in a direction chosen by the implementation, presumably the direction that is fastest. This includes the possibility that **Round** *x* is the same value as **RoundEven** *x* for all values of *x*. The operand *x* must be a scalar or vector whose component type is floating-point. *Result Type* and the type of *x* must be the same type. Results are computed per component.	
1	*<id>* *x*

RoundEven Result is the value equal to the nearest whole number to *x*. A fractional part of 0.5 will round toward the nearest even whole number. (Both 3.5 and 4.5 for *x* will be 4.0.) The operand *x* must be a scalar or vector whose component type is floating-point. *Result Type* and the type of *x* must be the same type. Results are computed per component.	
2	*<id>* *x*

Trunc

Result is the value equal to the nearest whole number to x whose absolute value is not larger than the absolute value of x.

The operand x must be a scalar or vector whose component type is floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

3	*<id>*
	x

FAbs

Result is x if $x \geq 0$; otherwise result is $-x$.

The operand x must be a scalar or vector whose component type is floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

4	*<id>*
	x

SAbs

Result is x if $x \geq 0$; otherwise result is $-x$, where x is interpreted as a signed integer.

Result Type and the type of x must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

5	*<id>*
	x

FSign

Result is 1.0 if $x > 0$, 0.0 if $x = 0$, or -1.0 if $x < 0$.

The operand x must be a scalar or vector whose component type is floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

6	*<id>*
	x

SSign

Result is 1 if $x > 0$, 0 if $x = 0$, or -1 if $x < 0$, where x is interpreted as a signed integer.

Result Type and the type of x must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

7	*<id>*
	x

Floor	
Result is the value equal to the nearest whole number that is less than or equal to *x*.	
The operand *x* must be a scalar or vector whose component type is floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	
8	*<id>*
	x

Ceil	
Result is the value equal to the nearest whole number that is greater than or equal to *x*.	
The operand *x* must be a scalar or vector whose component type is floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	
9	*<id>*
	x

Fract	
Result is *x* - **floor** *x*.	
The operand *x* must be a scalar or vector whose component type is floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	
10	*<id>*
	x

Radians	
Converts *degrees* to radians, i.e., *degrees* * π / 180.	
The operand *degrees* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.	
Result Type and the type of *degrees* must be the same type. Results are computed per component.	
11	*<id>*
	degrees

Degrees	
Converts *radians* to degrees, i.e., *radians* * 180 / π.	
The operand *radians* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.	
Result Type and the type of *radians* must be the same type. Results are computed per component.	
12	*<id>*
	radians

Sin	
The standard trigonometric sine of *x* radians.	
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	
13	*<id>*
	x

Cos	
The standard trigonometric cosine of *x* radians.	
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	
14	*<id>*
	x

Tan	
The standard trigonometric tangent of *x* radians.	
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	
15	*<id>*
	x

Asin	
Arc sine. Result is an angle, in radians, whose sine is *x*. The range of result values is [-π / 2, π / 2]. Result is undefined if **abs** *x* > 1.	
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	
16	*<id>*
	x

Acos	
Arc cosine. Result is an angle, in radians, whose cosine is *x*. The range of result values is [0, π]. Result is undefined if **abs** *x* > 1.	
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.	
Result Type and the type of *x* must be the same type. Results are computed per component.	

17	*<id>* *x*

Atan

Arc tangent. Result is an angle, in radians, whose tangent is *y_over_x*. The range of result values is [-π, π].

The operand *y_over_x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of *y_over_x* must be the same type. Results are computed per component.

18	*<id>* *y_over_x*

Sinh

Hyperbolic sine of *x* radians.

The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of *x* must be the same type. Results are computed per component.

19	*<id>* *x*

Cosh

Hyperbolic cosine of *x* radians.

The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of *x* must be the same type. Results are computed per component.

20	*<id>* *x*

Tanh

Hyperbolic tangent of *x* radians.

The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of *x* must be the same type. Results are computed per component.

21	*<id>* *x*

Asinh		
Arc hyperbolic sine; result is the inverse of **sinh**.		
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.		
Result Type and the type of *x* must be the same type. Results are computed per component.		
22	*<id>* *x*	

Acosh		
Arc hyperbolic cosine; Result is the non-negative inverse of **cosh**. Result is undefined if *x* < 1.		
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.		
Result Type and the type of *x* must be the same type. Results are computed per component.		
23	*<id>* *x*	

Atanh		
Arc hyperbolic tangent; result is the inverse of tanh. Result is undefined if **abs** $x \geq 1$.		
The operand *x* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.		
Result Type and the type of *x* must be the same type. Results are computed per component.		
24	*<id>* *x*	

Atan2		
Arc tangent. Result is an angle, in radians, whose tangent is *y* / *x*. The signs of *x* and *y* are used to determine what quadrant the angle is in. The range of result values is [-π, π] . Result is undefined if *x* and *y* are both 0.		
The operand *x* and *y* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.		
Result Type and the type of all operands must be the same type. Results are computed per component.		
25	*<id>* *y*	*<id>* *x*

Pow		
Result is *x* raised to the *y* power; x^y. Result is undefined if *x* < 0. Result is undefined if *x* = 0 and $y \leq 0$.		
The operand *x* and *y* must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.		
Result Type and the type of all operands must be the same type. Results are computed per component.		
26	*<id>* *x*	*<id>* *y*

Exp

Result is the natural exponentiation of x; e^x.

The operand x must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

27	*<id>*
	x

Log

Result is the natural logarithm of x, i.e., the value y which satisfies the equation $x = e^y$. Result is undefined if $x \leq 0$.

The operand x must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

28	*<id>*
	x

Exp2

Result is 2 raised to the x power; 2^x.

The operand x must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

29	*<id>*
	x

Log2

Result is the base-2 logarithm of x, i.e., the value y which satisfies the equation $x = 2^y$. Result is undefined if $x \leq 0$.

The operand x must be a scalar or vector whose component type is 16-bit or 32-bit floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

30	*<id>*
	x

Sqrt

Result is the square root of x. Result is undefined if $x < 0$.

The operand x must be a scalar or vector whose component type is floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

31	*<id>*
	x

InverseSqrt

Result is the reciprocal of **sqrt** x. Result is undefined if $x \leq 0$.

The operand x must be a scalar or vector whose component type is floating-point.

Result Type and the type of x must be the same type. Results are computed per component.

32	*<id>*
	x

Determinant

Result is the determinant of x.

The operand x must be a square matrix.

Result Type must be the same type as the component type in the columns of x.

33	*<id>*
	x

MatrixInverse

Result is a matrix that is the inverse of x. The values in the result are undefined if x is singular or poorly conditioned (nearly singular).

The operand x must be a square matrix.

Result Type and the type of x must be the same type.

34	*<id>*
	x

Modf

Result is the fractional part of x and stores through i the whole number part (as a whole-number floating-point value). Both the result and the output parameter will have the same sign as x.

The operand x must be a scalar or vector whose component type is floating-point.

The operand i must have a pointer type.

Result Type, the type of x, and the type i points to must all be the same type. Results are computed per component.

35	*<id>*	*<id>*
	x	*i*

ModfStruct

Same semantics as in **Modf**, except that the entire result is in the instruction's result; there is not a pointer operand to write through.

Result Type must be an **OpTypeStruct** with two members. Member 0 holds the fractional part. Member 1 holds the whole number part. These two members, and the *Result Type* must all have the same type. This structure type must be explicitly declared by the module.

36	*<id>*
	x

FMin

Result is *y* if *y* < *x*; otherwise result is *x*. Which operand is the result is undefined if one of the operands is a NaN.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

37	*<id>*	*<id>*
	x	*y*

UMin

Result is *y* if *y* < *x*; otherwise result is *x*, where *x* and *y* are interpreted as unsigned integers.

Result Type and the type of *x* and *y* must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

38	*<id>*	*<id>*
	x	*y*

SMin

Result is *y* if *y* < *x*; otherwise result is *x*, where *x* and *y* are interpreted as signed integers.

Result Type and the type of *x* and *y* must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

39	*<id>*	*<id>*
	x	*y*

FMax

Result is *y* if *x* < *y*; otherwise result is *x*. Which operand is the result is undefined if one of the operands is a NaN.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

40	*<id>*	*<id>*
	x	*y*

UMax

Result is *y* if *x* < *y*; otherwise result is *x*, where *x* and *y* are interpreted as unsigned integers.

Result Type and the type of *x* and *y* must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

41	<id>	<id>
	x	*y*

SMax

Result is *y* if *x* < *y*; otherwise result is *x*, where *x* and *y* are interpreted as signed integers.

Result Type and the type of *x* and *y* must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

42	<id>	<id>
	x	*y*

FClamp

Result is min(max(*x*, *minVal*), *maxVal*). Result is undefined if *minVal* > *maxVal*.The semantics used by min() and max() are those of FMin and FMax.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

43	<id>	<id>	<id>
	x	*minVal*	*maxVal*

UClamp

Result is min(max(*x*, *minVal*), *maxVal*), where *x*, *minVal* and *maxVal* are interpreted as unsigned integers. Result is undefined if *minVal* > *maxVal*.

Result Type and the type of the operands must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

44	<id>	<id>	<id>
	x	*minVal*	*maxVal*

SClamp

Result is min(max(*x*, *minVal*), *maxVal*), where *x*, *minVal* and *maxVal* are interpreted as signed integers. Result is undefined if *minVal* > *maxVal*.

Result Type and the type of the operands must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

45	<id>	<id>	<id>
	x	*minVal*	*maxVal*

FMix

Result is the linear blend of *x* and *y*, i.e., $x * (1 - a) + y * a$.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

46	*<id>*	*<id>*	*<id>*
	x	*y*	*a*

Step

Result is 0.0 if $x < edge$; otherwise result is 1.0.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

48	*<id>*	*<id>*
	edge	*x*

SmoothStep

Result is 0.0 if $x \leq edge0$ and 1.0 if $x \geq edge1$ and performs smooth Hermite interpolation between 0 and 1 when $edge0 < x < edge1$. This is equivalent to:

$t * t * (3 - 2 * t)$, where $t = $ clamp $((x - edge0) / (edge1 - edge0), 0, 1)$

Result is undefined if $edge0 \geq edge1$.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

49	*<id>*	*<id>*	*<id>*
	edge0	*edge1*	*x*

Fma

Computes $a * b + c$. In uses where this operation is decorated with **NoContraction**:

- **fma** is considered a single operation, whereas the expression $a * b + c$ is considered two operations.

- The precision of **fma** can differ from the precision of the expression $a * b + c$.

- **fma** will be computed with the same precision as any other **fma** decorated with **NoContraction**, giving invariant results for the same input values of a, b, and c.

Otherwise, in the absence of a **NoContraction** decoration, there are no special constraints on the number of operations or difference in precision between **fma** and the expression $a * b + c$.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

50	*<id>*	*<id>*	*<id>*
	a	*b*	*c*

Frexp

Splits x into a floating-point significand in the range [0.5, 1.0) and an integral exponent of two, such that:

$$x = significand * 2^{exponent}$$

The *significand* is the result and the exponent is returned through the pointer-parameter *exp*. For a floating-point value of zero, the significand and exponent are both zero. For a floating-point value that is an infinity or is not a number, the result is undefined.

If an implementation supports negative 0, **Frexp** -0 should result in -0; otherwise it will result in 0.

The operand x must be a scalar or vector whose component type is floating-point.

The *exp* operand must be a pointer to a scalar or vector with integer component type, with 32-bit component width. The number of components in x and what *exp* points to must be the same.

Result Type must be the same type as the type of x. Results are computed per component.

51	*<id>*	*<id>*
	x	*exp*

FrexpStruct

Same semantics as in **Frexp**, except that the entire result is in the instruction's result; there is not a pointer operand to write through.

Result Type must be an **OpTypeStruct** with two members. Member 0 must have the same type as the type of x. Member 0 holds the *significand*. Member 1 must be a scalar or vector with integer component type, with 32-bit component width. Member 1 holds *exponent*. These two members must have the same number of components. This structure type must be explicitly declared by the module.

52	*<id>*
	x

Ldexp

Builds a floating-point number from *x* and the corresponding integral exponent of two in *exp*:

significand $* 2^{exponent}$

If this product is too large to be represented in the floating-point type, the result is undefined. If *exp* is greater than +128 (single precision) or +1024 (double precision), the result undefined. If *exp* is less than -126 (single precision) or -1022 (double precision), the result may be flushed to zero. Additionally, splitting the value into a significand and exponent using **frexp** and then reconstructing a floating-point value using **ldexp** should yield the original input for zero and all finite non-denormalized values.

The operand *x* must be a scalar or vector whose component type is floating-point.

The *exp* operand must be a scalar or vector with integer component type. The number of components in *x* and *exp* must be the same.

Result Type must be the same type as the type of *x*. Results are computed per component.

53	*<id>*	*<id>*
	x	*exp*

PackSnorm4x8

First, converts each component of the normalized floating-point value *v* into 8-bit integer values. These are then packed into the result.

The conversion for component *c* of *v* to fixed point is done as follows:

round(clamp(*c*, -1, +1) * 127.0)

The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.

The *v* operand must be a vector of 4 components whose type is a 32-bit floating-point.

Result Type must be a 32-bit integer type.

54	*<id>*
	v

PackUnorm4x8

First, converts each component of the normalized floating-point value *v* into 8-bit integer values. These are then packed into the result.

The conversion for component *c* of *v* to fixed point is done as follows:

round(clamp(*c*, 0, +1) * 255.0)

The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.

The *v* operand must be a vector of 4 components whose type is a 32-bit floating-point.

Result Type must be a 32-bit integer type.

| 55 | *<id>* |
| | *v* |

PackSnorm2x16

First, converts each component of the normalized floating-point value *v* into 16-bit integer values. These are then packed into the result.

The conversion for component *c* of *v* to fixed point is done as follows:

round(clamp(*c*, -1, +1) * 32767.0)

The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.

The *v* operand must be a vector of 2 components whose type is a 32-bit floating-point.

Result Type must be a 32-bit integer type.

| 56 | *<id>* |
| | *v* |

PackUnorm2x16

First, converts each component of the normalized floating-point value *v* into 16-bit integer values. These are then packed into the result.

The conversion for component *c* of *v* to fixed point is done as follows:

round(clamp(*c*, 0, +1) * 65535.0)

The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.

The *v* operand must be a vector of 2 components whose type is a 32-bit floating-point.

Result Type must be a 32-bit integer type.

57	*<id>*
	v

PackHalf2x16

Result is the unsigned integer obtained by converting the components of a two-component floating-point vector to the 16-bit **OpTypeFloat**, and then packing these two 16-bit integers into a 32-bit unsigned integer. The first vector component specifies the 16 least-significant bits of the result; the second component specifies the 16 most-significant bits.

The *v* operand must be a vector of 2 components whose type is a 32-bit floating-point.

Result Type must be a 32-bit integer type.

58	*<id>*
	v

PackDouble2x32

Result is the double-precision value obtained by packing the components of *v* into a 64-bit value. If an IEEE 754 Inf or NaN is created, it will not signal, and the resulting floating-point value is unspecified. Otherwise, the bit-level representation of *v* is preserved. The first vector component specifies the 32 least significant bits; the second component specifies the 32 most significant bits.

The *v* operand must be a vector of 2 components whose type is a 32-bit integer.

Result Type must be a 64-bit floating-point scalar.

Use of this instruction requires declaration of the **Float64** capability.

59	*<id>*
	v

UnpackSnorm2x16

First, unpacks a single 32-bit unsigned integer *p* into a pair of 16-bit signed integers. Then, each component is converted to a normalized floating-point value to generate the result. The conversion for unpacked fixed-point value *f* to floating point is done as follows:

clamp(*f* / 32767.0, -1, +1)

The first component of the result will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.

The *p* operand must be a scalar with 32-bit integer type.

Result Type must be a vector of 2 components whose type is 32-bit floating point.

60	*<id>*
	p

UnpackUnorm2x16

First, unpacks a single 32-bit unsigned integer p into a pair of 16-bit unsigned integers. Then, each component is converted to a normalized floating-point value to generate the result. The conversion for unpacked fixed-point value f to floating point is done as follows:

f / 65535.0

The first component of the result will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.

The p operand must be a scalar with 32-bit integer type.

Result Type must be a vector of 2 components whose type is 32-bit floating point.

| 61 | *<id>* |
| | p |

UnpackHalf2x16

Result is the two-component floating-point vector with components obtained by unpacking a 32-bit unsigned integer into a pair of 16-bit values, interpreting those values as 16-bit floating-point numbers according to the OpenGL Specification, and converting them to 32-bit floating-point values. Subnormal numbers are either preserved or flushed to zero, consistently within an implemenation.

The first component of the vector is obtained from the 16 least-significant bits of v; the second component is obtained from the 16 most-significant bits of v.

The v operand must be a scalar with 32-bit integer type.

Result Type must be a vector of 2 components whose type is 32-bit floating point.

| 62 | *<id>* |
| | v |

UnpackSnorm4x8

First, unpacks a single 32-bit unsigned integer p into four 8-bit signed integers. Then, each component is converted to a normalized floating-point value to generate the result. The conversion for unpacked fixed-point value f to floating point is done as follows:

clamp(f / 127.0, -1, +1)

The first component of the result will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.

The p operand must be a scalar with 32-bit integer type.

Result Type must be a vector of 4 components whose type is 32-bit floating point.

| 63 | *<id>* |
| | p |

UnpackUnorm4x8

First, unpacks a single 32-bit unsigned integer *p* into four 8-bit unsigned integers. Then, each component is converted to a normalized floating-point value to generate the result. The conversion for unpacked fixed-point value *f* to floating point is done as follows:

f / 255.0

The first component of the result will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.

The *p* operand must be a scalar with 32-bit integer type.

Result Type must be a vector of 4 components whose type is 32-bit floating point.

64	*<id>*
	p

UnpackDouble2x32

Result is the two-component unsigned integer vector representation of *v*. The bit-level representation of *v* is preserved. The first component of the vector contains the 32 least significant bits of the double; the second component consists of the 32 most significant bits.

The *v* operand must be a scalar whose type is 64-bit floating point.

Result Type must be a vector of 2 components whose type is a 32-bit integer.

Use of this instruction requires declaration of the **Float64** capability.

65	*<id>*
	v

Length

Result is the length of vector *x*, i.e., $\mathrm{sqrt}(x[0]^2 + x[1]^2 + \dots)$.

The operand *x* must be a scalar or vector whose component type is floating-point.

Result Type must be a scalar of the same type as the component type of *x*.

66	*<id>*
	x

Distance

Result is the distance between *p0* and *p1*, i.e., length(*p0* - *p1*).

The operands must all be a scalar or vector whose component type is floating-point.

Result Type must be a scalar of the same type as the component type of the operands.

67	*<id>*	*<id>*
	p0	*p1*

Cross

Result is the cross product of *x* and *y*, i.e., the resulting components are, in order:

x[1] * *y*[2] - *y*[1] * *x*[2]

x[2] * *y*[0] - *y*[2] * *x*[0]

x[0] * *y*[1] - *y*[0] * *x*[1]

All the operands must be vectors of 3 components of a floating-point type.

Result Type and the type of all operands must be the same type.

68	*<id>*	*<id>*
	x	*y*

Normalize

Result is the vector in the same direction as *x* but with a length of 1.

The operand *x* must be a scalar or vector whose component type is floating-point.

Result Type and the type of *x* must be the same type.

69	*<id>*
	x

FaceForward

If the dot product of *Nref* and *I* is negative, the result is *N*, otherwise it is *-N*.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type.

70	*<id>*	*<id>*	*<id>*
	N	*I*	*Nref*

Reflect

For the incident vector *I* and surface orientation *N*, the result is the reflection direction:

I - 2 * dot(*N*, *I*) * *N*

N must already be normalized in order to achieve the desired result.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type.

71	*<id>*	*<id>*
	I	*N*

Refract

For the incident vector *I* and surface normal *N*, and the ratio of indices of refraction *eta*, the result is the refraction vector. The result is computed by

k = 1.0 - *eta* * *eta* * (1.0 - dot(*N*, *I*) * dot(*N*, *I*))

if k < 0.0 the result is 0.0

otherwise, the result is *eta* * *I* - (*eta* * dot(*N*, *I*) + sqrt(k)) * *N*

The input parameters for the incident vector *I* and the surface normal *N* must already be normalized to get the desired results.

The type of *I* and *N* must be a scalar or vector with a floating-point component type.

The type of *eta* must be a 16-bit or 32-bit floating-point scalar.

Result Type, the type of *I*, and the type of *N* must all be the same type.

72	*<id>*	*<id>*	*<id>*
	I	*N*	*eta*

FindILsb

Integer least-significant bit.

Results in the bit number of the least-significant 1-bit in the binary representation of *Value*. If *Value* is 0, the result is -1.

Result Type and the type of *Value* must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

73	*<id>*
	Value

FindSMsb

Signed-integer most-significant bit, with *Value* interpreted as a signed integer.

For positive numbers, the result will be the bit number of the most significant 1-bit. For negative numbers, the result will be the bit number of the most significant 0-bit. For a *Value* of 0 or -1, the result is -1.

Result Type and the type of *Value* must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

This instruction is currently limited to 32-bit width components.

74	*<id>*
	Value

FindUMsb

Unsigned-integer most-significant bit.

Results in the bit number of the most-significant 1-bit in the binary representation of *Value*. If *Value* is 0, the result is -1.

Result Type and the type of *Value* must both be integer scalar or integer vector types. *Result Type* and operand types must have the same number of components with the same component width. Results are computed per component.

This instruction is currently limited to 32-bit width components.

| 75 | *<id>* |
| | *Value* |

InterpolateAtCentroid

Result is the value of the input *interpolant* sampled at a location inside both the pixel and the primitive being processed. The value obtained would be the same value assigned to the input variable if it were decorated as **Centroid**.

The operand *interpolant* must be a pointer to the **Input** Storage Class.

The operand *interpolant* must be a pointer to a scalar or vector whose component type is 32-bit floating-point.

This instruction is only valid in the **Fragment** execution model.

Result Type and the type that *interpolant* points to must be the same type.

Use of this instruction requires declaration of the **InterpolationFunction** capability.

| 76 | *<id>* |
| | *interpolant* |

InterpolateAtSample

Result is the value of the input *interpolant* variable at the location of sample number *sample*. If multisample buffers are not available, the input variable will be evaluated at the center of the pixel. If sample *sample* does not exist, the position used to interpolate the input variable is undefined.

The operand *interpolant* must be a pointer to the **Input** Storage Class.

The operand *interpolant* must be a pointer to a scalar or vector whose component type is 32-bit floating-point.

This instruction is only valid in the **Fragment** execution model.

The *sample* operand must be a scalar 32-bit integer.

Result Type and the type that *interpolant* points to must be the same type.

Use of this instruction requires declaration of the **InterpolationFunction** capability.

| 77 | *<id>* | *<id>* |
| | *interpolant* | *sample* |

InterpolateAtOffset

Result is the value of the input *interpolant* variable sampled at an offset from the center of the pixel specified by *offset*. The two floating-point components of *offset*, give the offset in pixels in the *x* and *y* directions, respectively. An *offset* of (0, 0) identifies the center of the pixel. The range and granularity of offsets supported are implementation-dependent.

The operand *interpolant* must be a pointer to the **Input** Storage Class.

The operand *interpolant* must be a pointer to a scalar or vector whose component type is 32-bit floating-point.

This instruction is only valid in the **Fragment** execution model.

The *offset* operand must be a vector of 2 components of 32-bit floating-point type.

Result Type and the type that *interpolant* points to must be the same type.

Use of this instruction requires declaration of the **InterpolationFunction** capability.

78	*<id>*	*<id>*
	interpolant	*offset*

NMin

Result is *y* if *y* < *x*; otherwise result is *x*. If one operand is a NaN, the other operand is the result. If both operands are NaN, the result is a NaN.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

79	*<id>*	*<id>*
	x	*y*

NMax

Result is *y* if *x* < *y*; otherwise result is *x*. If one operand is a NaN, the other operand is the result. If both operands are NaN, the result is a NaN.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

80	*<id>*	*<id>*
	x	*y*

NClamp

Result is min(max(*x*, *minVal*), *maxVal*). Result is undefined if *minVal* > *maxVal*.The semantics used by min() and max() are those of NMin and NMax.

The operands must all be a scalar or vector whose component type is floating-point.

Result Type and the type of all operands must be the same type. Results are computed per component.

81	<id>	<id>	<id>
	x	minVal	maxVal

A   Changes

A.1   Changes from Version 0.99, Revision 1

- Fork the revision stream, changes section, etc. from the core specification, so this specification has its own, starting numbering at revision 1. This document now lives independently.
- Added integer versions of abs, sign, min, max, and clamp.
- Removed floatBitsToInt, floatBitsToUint, intBitsToFloat, and uintBitsToFloat; these can be handled with **OpBitcast**.
- Removed fTransform, not needed.
- Fixed internal bugs

  – 13721: Add **OpTypeStruct**-result versions of **Modf** and **Frexp**: **ModfStruct** and **FrexpStruct**.

- Fixed public bugs

  – 1322: GLSL.std.450 frexp wasn't saying the *exp* argument was a pointer to the result

A.2   Changes from Version 0.99, Revision 2

- Moved AddCarry, SubBorrow, and MulExtended type of instructions to the core specification.
- Added integer variant of **Mix**, creating **FMix** and **IMix** (14480).
- Modified spellings to be more regular (14614).

A.3   Changes from Version 0.99, Revision 3

- Add "N" version of **Min**, **Max**, and **Clamp**, creating a version that favors non-NaN operands over NaN operands.
- Bug 15452 Remove **IMix**.
- Bug 15300 Be more consistent that the **InterpolateAt** instructions take a pointer.
- Bug 14548 Document the **Capability** needed for **Double2x32** and **InterpolateAt** instructions.

A.4   Changes from Version 1.00, Revision 1

- Bug 14548 Document the **Capability** needed for **UnpackDouble2x32**.

A.5   Changes from Version 1.00, Revision 2

- Change **precise** to **NoContraction**

A.6   Changes from Version 1.00, Revision 3

- Allow both 16-bit and 32-bit floating-point types in most places where before only 32-bit floating-point types were allowed. This does not effect whether 16-bit floating point types are allowed, which is selected independently. Since 16-bit types were historically disallowed, this is a backward compatible change.
- Fix Khronos internal issue #109: be more clear for NMin/NMax: If both operands are NaN, the result is a NaN.

A.7   Changes from Version 1.00, Revision 4

• Be clear about **UnpackHalf2x16** denorm rules.