OpenCL 1.2 Extended Instruction Set Specification (Provisional)

Boaz Ouriel, Intel

Version 0.99, Revision 30

April 2, 2015

REVISION HISTORY

NUMBER	DATE	DESCRIPTION	NAME
1	Aug 2014	Created	jk
29	Mar 2015	Provisional Release	jk
30	2-Apr-2015	Provisional Release	jk

Contents

1 Introduction

This is the specification of **OpenCL.std.12** extended instruction set.

The library is imported into a SPIR-V module in the following manner:

<ext-inst-id> OpExtInstImport "OpenCL.std.12"

The library can only be imported when **Memory Model** is set to **OpenCL12**

2 Binary Form

This section contains the semantics and exact form of execution of OpenCL 1.2 extended instructions using the **OpExtInst** instruction.

In this section we use the following naming conventions:

- *void* denote an **OpTypeVoid**.

- *half*, *float* and *double* denote an **OpTypeFloat** with a width of 16, 32 and 64 bits respectively.

- *i8*, *i16*, *i32* and *i64* denote an **OpTypeInt** with a width of 8, 16, 32 and 64 bits respectively.

- *bool* denotes an **OpTypeBool**.

- *size_t* denotes an i32 when the **Addressing Model** is **Physical32** and i64 when the **Addressing Model** is **Physical64**.

- *vector(n)* denotes an **OpTypeVector** where *n* indicates the component count.

  – *vector(n₁, n₂, ..., nᵢ)* abbreviates *vector(n₁)*, *vector(n₂)*, ... or *vector(nᵢ)*.

- *integer* denotes *i8*, *i16*, *i32* or *i64*.

- *floating-point* denotes *half*, *float*, *double*.

- *pointer*(*storage*) denotes an **OpTypePointer** which points to *storage* **Storage Class**.

  – *pointer*(*constant*) denotes an OpTypePointer with **UniformConstant Storage Class**.

  – *pointer*(*global*) denotes an OpTypePointer with **WorkgroupGlobal Storage Class**.

  – *pointer*(*local*) denotes an OpTypePointer with **WorkgroupLocal Storage Class**.

  – *pointer*(*private*) denotes an OpTypePointer with **Private Storage Class**.

  – *pointer(s₁, s₂, ..., sᵢ)* abbreviates *pointer(s₁)*, *pointer(s₂)*, ... or *pointer(sᵢ)*.

- *image* defines all types of image memory objects (See image encoding section).

- *sampler* a SPIR-V sampler object (See sampler encoding section).

2.1   Math extended instructions

This section describes the list of external math instructions. The external math instructions are categorized into the following:

- A list of instructions that have scalar or vector argument versions, and,

- A list of instructions that only take scalar float arguments.

The vector versions of the math instructions operate component-wise. The description is per-component.

The math instructions are not affected by the prevailing rounding mode in the calling environment, and always return the same value as they would if called with the round to nearest even rounding mode.

acos						
Compute the arc cosine of *x*.						
Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.						
All of the operands, including the *Result Type* operand, must be of the same type.						
6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	0	*<id>* *x*

acosh

Compute the inverse hyperbolic cosine of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	1	*<id>* *x*

acospi

Compute *acos(x)* / π.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	2	*<id>* *x*

asin

Compute the arc sine of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	3	*<id>* *x*

asinh

Compute the inverse hyperbolic sine of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	4	*<id>* *x*

asinpi

Compute *asin*(*x*) / π.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	<id> Result Type	Result <id>	extended instructions set <id>	5	<id> x

atan

Compute the arc tangent of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	<id> Result Type	Result <id>	extended instructions set <id>	6	<id> x

atan2

Compute the arc tangent of *y* / *x*.

Result Type,*y* and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	<id> Result Type	Result <id>	extended instructions set <id>	7	<id> y	<id> x

atanh

Compute the hyperbolic arc tangent of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	<id> Result Type	Result <id>	extended instructions set <id>	8	<id> x

atanpi

Compute *atan*(*x*) / π.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	9	*<id>* *x*

atan2pi

Compute *atan2*(*y*, *x*) / π.

Result Type,*y* and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	10	*<id>* *y*	*<id>* *x*

cbrt

Compute the cube-root of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	11	*<id>* *x*

ceil

Round *x* to integral value using the round to positive infinity rounding mode.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	12	*<id>* *x*

copysign

Returns *x* with its sign changed to match the sign of *y*.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	<id> Result Type	Result <id>	extended instructions set <id>	13	<id> x	<id> y

cos

Compute the cosine of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	<id> Result Type	Result <id>	extended instructions set <id>	14	<id> x

cosh

Compute the hyperbolic cosine of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	<id> Result Type	Result <id>	extended instructions set <id>	15	<id> x

cospi

Compute *cos(x)* / π.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	<id> Result Type	Result <id>	extended instructions set <id>	16	<id> x

erfc

Complementary error function of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	17	*<id>* *x*

erf

Error function of *x* encountered in integrating the normal distribution.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	18	*<id>* *x*

exp

Compute the base-e exponential of *x*. (i.e. e^x)

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	19	*<id>* *x*

exp2

Computes 2 raised to the power of *x*. (i.e. 2^x)

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	20	*<id>* *x*

exp10

Computes 10 raised to the power of *x*. (i.e. 10^x)

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	21	*<id>* *x*

expm1

Computes $e^x - 1.0$.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	22	*<id>* *x*

fabs

Compute the absolute value of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	23	*<id>* *x*

fdim

Compute *x* - *y* if *x* > *y*, *+0* if *x* is less than or equal to *y*.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	24	*<id>* *x*	*<id>* *y*

floor

Round *x* to the integral value using the round to negative infinity rounding mode.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	25	*<id>* *x*

fma

Compute the correctly rounded floating-point representation of the sum of *c* with the infinitely precise product of *a* and *b*.Rounding of intermediate products shall not occur. Edge case behavior is per the IEEE 754-2008 standard.

Result Type,*a*,*b* and *c* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	26	*<id>* *a*	*<id>* *b*	*<id>* *c*

fmax

Returns *y* if *x < y*, otherwise it returns *x*. If one argument is a NaN, *Fmax* returns the other argument. If both arguments are NaNs, *Fmax* returns a NaN.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

Note: fmax behave as defined by C99 and may not match the IEEE 754-2008 definition for maxNum with regard to signaling NaNs.Specifically, signaling NaNs may behave as quiet NaNs

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	27	*<id>* *x*	*<id>* *y*

fmin

Returns *y* if *y < x*, otherwise it returns *x*. If one argument is a NaN, *Fmin* returns the other argument. If both arguments are NaNs, *Fmin* returns a NaN.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

Note: fmin behave as defined by C99 and may not match the IEEE 754-2008 definition for minNum with regard to signaling NaNs.Specifically, signaling NaNs may behave as quiet NaNs

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	28	*<id>* *x*	*<id>* *y*

fmod

Modulus. Returns *x - y * trunc (x/y)*.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	29	*<id>* *x*	*<id>* *y*

fract

Returns *fmin(x - floor(x)*, 0x1.fffffep-1f. *floor(x)* is returned in *ptr*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

ptr must be a *pointer*(*global*, *local*, *private*) to *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type, or must be a pointer to the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	30	*<id>* *x*	*<id>* *ptr*

frexp

Extract the mantissa and exponent from *x*. The *Result Type* holds the mantissa, and *exp* points to the exponent. For each component the mantissa returned is a *floating-point* with magnitude in the interval [1/2, 1) or 0. Each component of *x* equals mantissa returned $* 2^{exp}$.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

exp must be a *pointer*(*global*, *local*, *private*) to *i32* or *vector(2,3,4,8,16)* of *i32* values.

Result Type and *x* operands must be of the same type. *exp* operand must point to an *i32* with the same component count as *Result Type* and *x* operands.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	31	*<id>* *x*	*<id>* *exp*

hypot

Compute the value of the square root of $x^2 + y^2$ without undue overflow or underflow.

Result Type, *x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	32	*<id>* *x*	*<id>* *y*

ilogb

Return the exponent of *x* as an *i32* value.

Result Type must be *i32* or *vector(2,3,4,8,16)* of *i32* values.

x must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

Result Type and *x* operands must have the same component count.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	33	*<id>* *x*

ldexp

Multiply *x* by 2 to the power *k*.

k must be *i32* or *vector(2,3,4,8,16)* of *i32* values.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

Result Type and *x* operands must be of the same type. *exp* operand must have the same component count as *Result Type* and *x* operands.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	34	*<id>* *x*	*<id>* *k*

lgamma

Log gamma function of *x*. Returns the natural logarithm of the absolute value of the gamma function.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	35	*<id>* *x*

lgamma_r

Log gamma function of *x*. Returns the natural logarithm of the absolute value of the gamma function. The sign of the gamma function is returned in the *signp* operand

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

singp must be a *pointer*(*global*, *local*, *private*) to *i32* or *vector(2,3,4,8,16)* of *i32* values.

Result Type and *x* operands must be of the same type. *singp* operand must point to an *i32* with the same component count as *Result Type* and *x* operands.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	36	*<id>* *x*	*<id>* *singp*

log

Compute natural logarithm of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	37	*<id>* *x*

log2

Compute a base 2 logarithm of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	38	*<id>* *x*

log10

Compute a base 10 logarithm of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	39	*<id>* *x*

log1p

Compute $\log_e(1.0 + x)$.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	40	*<id>* *x*

logb

Compute the exponent of *x*, which is the integral part of $\log_r |x|$.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	41	*<id>* *x*

mad

mad approximates *a * b + c*. Whether or how the product of *a * b* is rounded and how supernormal or subnormal intermediate products are handled is not defined. mad is intended to be used where speed is preferred over accuracy

Result Type,a,b and *c* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

Note: For some usages, e.g.mad(*a*, *b*, *-a*b*), the definition of mad() is loose enough that almost any result is allowed from mad() for some values of *a* and *b*.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	42	*<id>* *a*	*<id>* *b*	*<id>* *c*

maxmag

Returns *x* if $|x| > |y|$, *y* if $|y| > |x|$, otherwise *fmax(x, y)*.

Result Type,x and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	43	*<id>* *x*	*<id>* *y*

minmag

Returns x if $|x| < |y|$, y if $|y| < |x|$, otherwise *fmin(x, y)*.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	44	*<id>* *x*	*<id>* *y*

modf

Decompose a *floating-point* number. The modf function breaks the argument *x* into integral and fractional parts, each of which has the same sign as the argument. It stores the integral part in the object pointed to by *iptr*

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

iptr must be a *pointer*(*global*, *local*, *private*) to *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type, or must be a pointer to the same type.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	45	*<id>* *x*	*<id>* *iptr*

nan

Returns a quiet NaN. The *nancode* may be placed in the significand of the resulting NaN.

nancode must be *i32* or *vector(2,3,4,8,16)* of *i32* values.

Result Type must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

Result Type and *nancode* operands must have the same component count.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	46	*<id>* *nancode*

nextafter

Computes the next representable *floating-point* value following *x* in the direction of *y*. Thus, if *y* is less than *x*, *nextafter*() returns the largest representable floating-point number less than *x*.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	47	*<id>* *x*	*<id>* *y*

---

pow

Compute *x* to the power *y*.

Result Type,*x*,*y* and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	48	*<id>* *x*	*<id>* *y*	*<id>* *x*

---

pown

Compute *x* to the power *y*, where *y* is an *i32* integer.

y must be *i32* or *vector(2,3,4,8,16)* of *i32* values.

Result Type must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

Result Type and *x* operands must be of the same type. *y* operand must have the same component count as *Result Type* and *x* operands.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	49	*<id>* *y*

---

powr

Compute *x* to the power *y*, where *y* is an integer.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	50	*<id>* *x*	*<id>* *y*

---

remainder

Compute the value r such that r = *x* - n*y, where n is the integer nearest the exact value of *x/y*. If there are two integers closest to *x/y*, n shall be the even one. If r is zero, it is given the same sign as *x*.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	51	*<id>* *x*	*<id>* *y*

---

remquo

The remquo function computes the value r such that r = *x* - k*y, where k is the integer nearest the exact value of *x/y*. If there are two integers closest to *x/y*, k shall be the even one. If r is zero, it is given the same sign as *x*. This is the same value that is returned by the *remainder* function. remquo also calculates the lower seven bits of the integral quotient *x/y*, and gives that value the same sign as *x/y*. It stores this signed value in the object pointed to by *quo*.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

quo must be a *pointer*(*global*, *local*, *private*) to *i32* or *vector(2,3,4,8,16)* of *i32* values.

Result Type, *x* and *y* operands must be of the same type. *quo* operand must point to an *i32* with the same component count as *Result Type*, *x* and *y* operands.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	52	*<id>* *x*	*<id>* *y*	*<id>* *quo*

---

rint

Round *x* to integral value (using round to nearest even rounding mode) in floating-point format.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	53	*<id>* *x*

---

rootn

Compute *x* to the power 1/*y*.

y must be *i32* or *vector(2,3,4,8,16)* of *i32* values.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

Result Type and *x* operands must be of the same type. *y* operand must have the same component count as *Result Type* and *x* operands.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	54	*<id>* *x*	*<id>* *y*

round

Return the integral value nearest to *x* rounding halfway cases away from zero, regardless of the current rounding direction.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	55	*<id>* *x*

rsqrt

Compute inverse square root of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	56	*<id>* *x*

sin

Compute sine of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	57	*<id>* *x*

sincos

Compute sine and cosine of *x*. The computed sine is the return value and computed cosine is returned in *cosval*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

cosval must be a *pointer*(*global*, *local*, *private*) to *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type, or must be a pointer to the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	58	*<id>* *x*	*<id>* *cosval*

sinh

Compute hyperbolic sine of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	59	*<id>* *x*

sinpi

Compute *sin* (π x).

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	60	*<id>* *x*

sqrt

Compute square root of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	61	*<id>* *x*

tan

Compute tangent of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	62	*<id>* *x*

tanh

Compute hyperbolic tangent of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	63	*<id>* *x*

tanpi

Compute *tan* (π x).

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	64	*<id>* *x*

tgamma

Compute the gamma function of *x*.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	65	*<id>* *x*

trunc

Round *x* to integral value using the round to zero rounding mode.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	66	*<id>* *x*

half_cos

Compute cosine of x, where x must be in the range $-2^{16} \ldots +2^{16}$.

Result Type and x must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	67	*<id>* *x*

half_divide

Compute x / y.

Result Type,x and y must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	68	*<id>* *x*	*<id>* *y*

half_exp

Compute the base-e exponential of x.

Result Type and x must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	69	*<id>* *x*

half_exp2

Compute the base- 2 exponential of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	70	*<id>* *x*

half_exp10

Compute the base- 10 exponential of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	71	*<id>* *x*

half_log

Compute natural logarithm of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	72	*<id>* *x*

half_log2

Compute a base 2 logarithm of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* Result Type	Result *<id>*	extended instructions set *<id>*	73	*<id>* x

half_log10

Compute a base 10 logarithm of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* Result Type	Result *<id>*	extended instructions set *<id>*	74	*<id>* x

half_powr

Compute *x* to the power *y*, where *x* is >= 0.

Result Type,*x* and *y* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

7	44	*<id>* Result Type	Result *<id>*	extended instructions set *<id>*	75	*<id>* x	*<id>* y

half_recip

Compute reciprocal of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	76	*<id>* *x*

half_rsqrt

Compute inverse square root of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	77	*<id>* *x*

half_sin

Compute sine of *x*, where *x* must be in the range $-2^{16} \dots +2^{16}$.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	78	*<id>* *x*

half_sqrt

Compute the square root of *x*.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	79	*<id>* *x*

half_tan

Compute tangent value of *x*, where *x* must be in the range $-2^{16} \ldots +2^{16}$.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

This function is implemented with a minimum of 10-bits of accuracy i.e. an ULP value $\Leftarrow$ 8192 ulp.

The support for denormal values is optional and may return any result allowed even when -cl-denormals-are-zero flag is not in force.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	80	*<id>* *x*

native_cos

Compute cosine of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	81	*<id>* *x*

native_divide

Compute *x / y* over an implementation-defined range. The maximum error is implementation-defined.

Result Type,*x* and *y* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	82	*<id>* *x*	*<id>* *y*

native_exp

Compute the base-e exponential of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	83	*<id>* *x*

native_exp2

Compute the base- 2 exponential of *x* over an implementation-defined range. The maximum error is implementation-defined..

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	84	*<id>* *x*

native_exp10

Compute the base- 10 exponential of *x* over an implementation-defined range. The maximum error is implementation-defined..

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	85	*<id>* *x*

native_log

Compute natural logarithm of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	86	*<id>* *x*

native_log2

Compute a base 2 logarithm of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	87	*<id>* *x*

native_log10

Compute a base 10 logarithm of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	88	*<id>* *x*

native_powr

Compute *x* to the power *y*, where *x* is >= 0.

Result Type,*x* and *y* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	89	*<id>* *x*	*<id>* *y*

native_recip

Compute reciprocal of *x* over an implementation-defined range. The range of x and y are implementation-defined. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	90	*<id>* *x*

native_rsqrt

Compute inverse square root of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	91	*<id>* *x*

native_sin

Compute sine of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	92	*<id>* *x*

native_sqrt

Compute the square root of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	93	*<id>* *x*

---

native_tan

Compute tangent value of *x* over an implementation-defined range. The maximum error is implementation-defined.

Result Type and *x* must be *float* or *vector(2,3,4,8,16)* of *float* values.

All of the operands, including the *Result Type* operand, must be of the same type.

The function may map to one or more native device instructions and will typically have better performance compared to the non native corresponding functions. Support for denormal values is implementation-defined for this function

| 6 | 44 | *<id>*
Result Type | Result *<id>* | extended
instructions
set *<id>* | 94 | *<id>*
x |

2.2 Integer instructions

This section describes the list of integer instructions that take scalar or vector arguments. The vector versions of the integer functions operate component-wise. The description is per-component.

---

s_abs

Returns |*x*|, where *x* is treated as signed integer.

Result Type and *x* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

| 6 | 44 | *<id>*
Result Type | Result *<id>* | extended
instructions
set *<id>* | 141 | *<id>*
x |

---

s_abs_diff

Returns | *x* - *y* | without modulo overflow, where *x* and *y* are treated as signed integers.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

| 7 | 44 | *<id>*
Result Type | Result *<id>* | extended
instructions
set *<id>* | 142 | *<id>*
x | *<id>*
y |

s_add_sat

Returns the saturated value of $x + y$, where x and y are treated as signed integers.

Result Type,x and y must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	143	*<id>* *x*	*<id>* *y*

u_add_sat

Returns the saturated value of $x + y$, where x and y are treated as unsigned integers.

Result Type,x and y must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	144	*<id>* *x*	*<id>* *y*

s_hadd

Returns the value of $(x + y) >> 1$, where x and y are treated as signed integers. The intermediate sum does not modulo overflow.

Result Type,x and y must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	145	*<id>* *x*	*<id>* *y*

u_hadd

Returns the value of $(x + y) >> 1$, where x and y are treated as unsigned integers. The intermediate sum does not modulo overflow.

Result Type,x and y must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	146	*<id>* *x*	*<id>* *y*

s_rhadd

Returns the value of $(x + y + 1) >> 1$, where *x* and *y* are treated as signed integers. The intermediate sum does not modulo overflow.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	147	*<id>* *x*	*<id>* *y*

u_rhadd

Returns the value of $(x + y + 1) >> 1$, where *x* and *y* are treated as unsigned integers. The intermediate sum does not modulo overflow.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	148	*<id>* *x*	*<id>* *y*

s_clamp

Returns *s_min*(*s_max*(*x,minval*),*maxval*). Results are undefined if *minval* > *maxval*.

Result Type,*x,minval* and *maxval* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	149	*<id>* *x*	*<id>* *minval*	*<id>* *maxval*

u_clamp

Returns *u_min*(*u_max*(*x,minval*),*maxval*). Results are undefined if *minval* > *maxval*.

Result Type,*x,minval* and *maxval* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	150	*<id>* *x*	*<id>* *minval*	*<id>* *maxval*

clz

Returns the number of leading 0-bits in *x*, starting at the most significant bit position. If *x* is 0, returns the size in bits of the type of *x* or component type of *x*, if *x* is a vector.

Result Type and *x* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	151	*<id>* *x*

ctz

Returns the count of trailing 0-bits in *x*. If *x* is 0, returns the size in bits of the type of *x* or component type of *x*, if *x* is a vector.

Result Type and *x* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	152	*<id>* *x*

s_mad_hi

Returns *mul_hi(a, b)* + *c*, where *a,b* and *c* are treated as signed integers.

Result Type,a,b and *c* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	153	*<id>* *a*	*<id>* *b*	*<id>* *c*

s_max

Returns *y* if *x* < *y*, otherwise it returns *x*, where *x* and *y* are treated as signed integers.

Result Type,x and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	156	*<id>* *x*	*<id>* *y*

u_max

Returns *y* if *x* < *y*, otherwise it returns *x*, where *x* and *y* are treated as unsigned integers.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	157	*<id>* *x*	*<id>* *y*

s_min

Returns *y* if *y* < *x*, otherwise it returns *x*, where *x* and *y* are treated as signed integers.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	158	*<id>* *x*	*<id>* *y*

u_min

Returns *y* if *y* < *x*, otherwise it returns *x*, where *x* and *y* are treated as unsigned integers.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	159	*<id>* *x*	*<id>* *y*

s_mul_hi

Computes *x* * *y* and returns the high half of the product of *x* and *y*, where *x* and *y* are treated as signed integers.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	160	*<id>* *x*	*<id>* *y*

rotate

For each element in *v*, the bits are shifted left by the number of bits given by the corresponding element in *i*. Bits shifted off the left side of the element are shifted back in from the right.

Result Type,*v* and *i* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	161	*<id>* *v*	*<id>* *i*

s_sub_sat

Returns the saturated value of *x* - *y*, where *x* and *y* are treated as signed integers.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	162	*<id>* *x*	*<id>* *y*

u_sub_sat

Returns the saturated value of *x* - *y*, where *x* and *y* are treated as unsigned integers.

Result Type,*x* and *y* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	163	*<id>* *x*	*<id>* *y*

u_upsample

When *hi* and *lo* component type is i8:

Result = ((upcast. . . to i16)*hi* << 8) | *lo*

When *hi* and *lo* component type is i16:

Result = ((upcast. . . to i32)*hi* << 8) | *lo*

When *hi* and *lo* component i32:

Result = ((upcast. . . to i64)*hi* << 8) | *lo*

hi and *lo* are treated as unsigned integers.

hi and *lo* must be *i8*, *i16* or *i32* or *vector(2,3,4,8,16)* of *i8*, *i16* or *i32* values.

Result Type must be *i16*, *i32* or *i64* or *vector(2,3,4,8,16)* of *i16*, *i32* or *i64* values.

hi and *lo* operands must be of the same type. When *hi* and *lo* component type is i8, the *Result Type* component type must be i16. When *hi* and *lo* component type is i16, the *Result Type* component type must be i32. When *hi* and *lo* component type is i32, the *Result Type* component type must be i64. *Result Type* must have the same component count as *hi* and *lo* operands.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	164	*<id>* *hi*	*<id>* *lo*

s_upsample

When *hi* and *lo* component type is i8:

Result = ((upcast. . . to i16)*hi* << 8) | *lo*

When *hi* and *lo* component type is i16:

Result = ((upcast. . . to i32)*hi* << 8) | *lo*

When *hi* and *lo* component i32:

Result = ((upcast. . . to i64)*hi* << 8) | *lo*

hi and *lo* are treated as signed integers.

hi and *lo* must be *i8*, *i16* or *i32* or *vector(2,3,4,8,16)* of *i8*, *i16* or *i32* values.

Result Type must be *i16*, *i32* or *i64* or *vector(2,3,4,8,16)* of *i16*, *i32* or *i64* values.

hi and *lo* operands must be of the same type. When *hi* and *lo* component type is i8, the *Result Type* component type must be i16. When *hi* and *lo* component type is i16, the *Result Type* component type must be i32. When *hi* and *lo* component type is i32, the *Result Type* component type must be i64. *Result Type* must have the same component count as *hi* and *lo* operands.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	165	*<id>* *hi*	*<id>* *lo*

popcount

Returns the number of non-zero bits in *x*.

Result Type and *x* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	166	*<id>* *x*

s_mad24

Multipy two 24-bit integer values *x* and *y* and add the 32-bit integer result to the 32-bit integer *z*. Refer to definition of s_mul24 to see how the 24-bit integer multiplication is performed.

Result Type,x,y and *z* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	167	*<id>* *x*	*<id>* *y*	*<id>* *z*

u_mad24

Multipy two 24-bit integer values *x* and *y* and add the 32-bit integer result to the 32-bit integer *z*. Refer to definition of u_mul24 to see how the 24-bit integer multiplication is performed.

Result Type,x,y and *z* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	168	*<id>* *x*	*<id>* *y*	*<id>* *z*

s_mul24

Multiply two 24-bit integer values *x* and *y*, where *x* and *y* are treated as signed integers. *x* and *y* are 32-bit integers but only the low 24-bits are used to perform the multiplication. s_mul24 should only be used when values in *x* and *y* are in the range $[-2^{23}, 2^{23}-1]$. If *x* and *y* are not in this range, the multiplication result is implementation-defined.

Result Type,x and *y* must be *i32* or *vector(2,3,4,8,16)* of *i32* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	<id> Result Type	Result <id>	extended instructions set <id>	169	<id> x	<id> y

u_mul24

Multiply two 24-bit integer values x and y, where x and y are treated as unsigned integers. x and y are 32-bit integers but only the low 24-bits are used to perform the multiplication. u_mul24 should only be used when values in x and y are in the range $[0, 2^{24}-1]$. If x and y are not in this range, the multiplication result is implementation-defined.

Result Type,x and y must be *i32* or *vector(2,3,4,8,16)* of *i32* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	<id> Result Type	Result <id>	extended instructions set <id>	170	<id> x	<id> y

u_abs

Returns |x|, where x is treated as unsigned integer.

Result Type and x must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	<id> Result Type	Result <id>	extended instructions set <id>	201	<id> x

u_abs_diff

Returns | x - y | without modulo overflow, where x and y are treated as unsigned integers.

Result Type,x and y must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	<id> Result Type	Result <id>	extended instructions set <id>	202	<id> x	<id> y

u_mul_hi

Computes $x * y$ and returns the high half of the product of x and y, where x and y are treated as unsigned integers.

Result Type,x and y must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	<id> Result Type	Result <id>	extended instructions set <id>	203	<id> x	<id> y

u_mad_hi

Returns *mul_hi(a, b) + c*, where *a,b* and *c* are treated as unsigned integers.

Result Type,a,b and *c* must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	204	*<id>* *a*	*<id>* *b*	*<id>* *c*

2.3 Common instructions

This section describes the the list of common instructions that take scalar or vector arguments. The vector versions of the integer functions operate component-wise. The description is per-component. The common instructions are implemented using the round to nearest even rounding mode.

fclamp

Returns *fmin(fmax(x, minval), maxval)*. Results are undefined if *minval > maxval*.

Result Type,x,minval and *maxval* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	95	*<id>* *x*	*<id>* *minval*	*<id>* *maxval*

degrees

Converts *radians* to degrees, i.e. $(180 / \pi)$ * *radians*.

Result Type and *radians* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	96	*<id>* *radians*

fmax_common

Returns *y* if *x < y*, otherwise it returns *x*. If x or y are infinite or NaN, the return values are undefined.

Result Type,x and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	97	*<id>* *x*	*<id>* *y*

fmin_common

Returns *y* if *y* < *x*, otherwise it returns *x*. If x or y are infinite or NaN, the return values are undefined.

Result Type,*x* and *y* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	98	*<id>* *x*	*<id>* *y*

mix

Returns the linear blend of *x* & *y* implemented as:

x + (y - x) * a

Result Type,*x*,*y* and *a* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

Note: This function can be implemented using contractions such as mad or fma

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	99	*<id>* *x*	*<id>* *y*	*<id>* *a*

radians

Converts *degrees* to radians, i.e. $(\pi / 180) * degrees$.

Result Type and *degrees* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	100	*<id>* *degrees*

step

Returns 0.0 if *x* < *edge*, otherwise it returns 1.0.

Result Type,*edge* and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	101	*<id>* *edge*	*<id>* *x*

---

smoothstep

Returns 0.0 if $x \Leftarrow edge_0$ and 1.0 if $x \geq edge_1$ and performs smooth Hermite interpolation between 0 and 1, when $edge_0 <$ x $< edge_1$.

This is equivalent to :

t = $fclamp((x - edge_0) / (edge_1 - edge_0), 0, 1)$;

return t * t * (3 - 2 * t);

Results are undefined if $edge_0 \geq edge_1$ or if x, $edge_0$ or $edge_1$ is a NaN.

Result Type,$edge_0$,$edge_1$ and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

Note: This function can be implemented using contractions such as mad or fma

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	102	*<id>* *$edge_0$*	*<id>* *$edge_1$*	*<id>* *x*

---

sign

Returns 1.0 if $x > 0$, -0.0 if $x = -0.0$, +0.0 if $x = +0.0$, or -1.0 if $x < 0$. Returns 0.0 if x is a NaN.

Result Type and *x* must be *floating-point* or *vector(2,3,4,8,16)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	103	*<id>* *x*

2.4 Geometric instructions

This section describes the the list of geometric instructions. In this section *x,y,z* and *w* denote the first, second, third and fourth component respectively, of vectors with 3 and four components.The geometric instructions are implemented using the round to nearest even rounding mode.

Note: The geometric functions can be implemented using contractions such as mad or fma

cross

Returns the cross product of p_0.xyz and p_1.xyz.

When the vector component count is 4, the w component returned will be 0.0.

Result Type, p_0 and p_1 must be *vector(3,4)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	104	*<id>* p_0	*<id>* p_1

distance

Returns the distance between p_0 and p_1. This is calculated as *length*(p_0 - p_1).

Result Type must be *floating-point*.

p_0 and p_1 must be *floating-point* or *vector(2,3,4)* of *floating-point* values.

p_0 and p_1 operands must have the same type. *Result Type*, p_0 and p_1 operands must have the same component type

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	105	*<id>* p_0	*<id>* p_1

length

Return the length of vector p, i.e. *sqrt*($p.x^2 + p.y^2 + \ldots$)

Result Type must be *floating-point*.

p must be *vector(2,3,4)* of *floating-point* values.

Result Type and p operands must have the same component type

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	106	*<id>* p

normalize

Returns a vector in the same direction as p but with a length of 1.

Result Type and p must be *floating-point* or *vector(2,3,4)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	107	*<id>* p

fast_distance

Returns *fast_length(p_0 - p_1)*.

Result Type must be *floating-point*.

p_0 and p_1 must be *floating-point* or *vector(2,3,4)* of *floating-point* values.

p_0 and p_1 operands must have the same type. *Result Type*, p_0 and p_1 operands must have the same component type

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	108	*<id>* p_0	*<id>* p_1

fast_length

Return the length of vector *p* computed as: *half_sqrt*($p.x^2 + p.y^2 + \dots$)

Result Type must be *floating-point*.

p must be *vector(2,3,4)* of *floating-point* values.

Result Type and *p* operands must have the same component type

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	109	*<id>* *p*

fast_normalize

Returns a vector in the same direction as *p* but with a length of 1 computed as:

$p * half\_rsqrt(p.x^2 + p.y^2 \dots)$

The result shall be within 8192 ulps error from the infinitely precise result of:

if (*all*(*p* == 0.0f)) { result = *p*; }

else { result = *p* / *sqrt*($p.x^2 + p.y^2 + \dots$); }

with the following exceptions :

1) If the sum of squares is greater than FLT_MAX then the value of the floating-point values in the result vector are undefined.

2) If the sum of squares is less than FLT_MIN then the implementation may return back *p*.

3) If the device is in "denorms are flushed to zero" mode, individual operand elements with magnitude less than *sqrt*(FLT_MIN) may be flushed to zero before proceeding with the calculation.

Result Type and *p* must be *floating-point* or *vector(2,3,4)* of *floating-point* values.

All of the operands, including the *Result Type* operand, must be of the same type.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	110	*<id>* *p*

2.5 Relational instructions

This section describes the the list of relational instructions that take scalar or vector arguments. The vector versions of the integer functions operate component-wise. The description is per-component.

---

bitselect

Each bit of the result is the corresponding bit of *a* if the corresponding bit of *c* is 0. Otherwise it is the corresponding bit of *b*.

Result Type,*a*,*b* and *c* must be *floating-point* or *integer* or *vector(2,3,4,8,16)* of *floating-point* or *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	186	*<id>* *a*	*<id>* *b*	*<id>* *c*

---

select

Each bit of the result is the corresponding bit of *a* if the corresponding bit of *c* is 0. Otherwise it is the corresponding bit of *b*.

c must be *integer* or *vector(2,3,4,8,16)* of *integer* values.

Result Type,*a* and *b* must be *floating-point* or *integer* or *vector(2,3,4,8,16)* of *floating-point* or *integer* values.

Result Type, *a* and *b* must have the same type. *c* operand must have the same component count and component bit width as the rest of the operands.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	187	*<id>* *a*	*<id>* *b*	*<id>* *c*

2.6 Vector Data Load and Store instructions

This section describes the list of instructions that allow reading and writing of vector types from a pointer to memory.

vloadn

Return a vector value which is read from address (*p* + (*offset* * *n*)).

The address computed as (*p* + (*offset* * *n*)) must be 8-bit aligned if *p* points to i8 value; 16-bit aligned if *p* points to i16 or half value; 32-bit aligned if *p* points to i32 or float value; 64-bit aligned if *p* points to i64 or double value.

offset must be *size_t*.

p must be a *pointer*(*global*, *local*, *private*, *constant*) to *floating-point*, *integer*.

Result Type must be *vector(2,3,4,8,16)* of *floating-point* or *integer* values.

Result Type component count must be equal to *n* and its component type must be equal to the type pointed by *p*.

n must be 2,3,4,8 or 16.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	171	*<id>* *offset*	*<id>* *p*	Literal Number *n*

vstoren

Write *data* vector value to the address (*p* + (*offset* * compCountOf(*data*))), where compCountOf(*data*) is equal to the component count of the vector *data*.

The address computed as (*p* + (*offset* * compCountOf(*data*))) must be 8-bit aligned if *p* points to i8 value; 16-bit aligned if *p* points to i16 or half value; 32-bit aligned if *p* points to i32 or float value; 64-bit aligned if *p* points to i64 or double value.

offset must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*, *local*, *private*) to *floating-point*, *integer*.

data must be *vector(2,3,4,8,16)* of *floating-point* or *integer* values.

data component type must be equal to the type pointed by *p*.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	172	*<id>* *data*	*<id>* *offset*	*<id>* *p*

vload_half

Reads a half value from the address (*p* + (*offset*)) and converts it to a float return value. The address computed as (*p* + (*offset*)) must be 16-bit aligned.

Result Type must be *float*.

offset must be *size_t*.

p must be a *pointer*(*global*, *local*, *private*, *constant*) to *half*.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	173	*<id>* *offset*	*<id>* *p*	

vload_halfn

Reads a half vector value from the address (*p* + (*offset* * *n*)) and converts it to a float vector return value. The address computed as (*p* + (*offset* * *n*)) must be 16-bit aligned.

offset must be *size_t*.

p must be a *pointer*(*global*, *local*, *private*, *constant*) to *half*.

Result Type must be *vector(2,3,4,8,16)* of *float* values.

Result Type component count must be equal to *n*.

n must be 2,3,4,8 or 16.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	174	*<id>* *offset*	*<id>* *p*	Literal Number *n*

vstore_half

Converts *data* float or double value to a half value and then write the converted value to the address (*p* + *offset*). The address computed as (*p* + *offset*) must be 16-bit aligned.

This function uses the default rounding mode when converting *data* to a half value. The default rounding mode is round to nearest even.

data must be *float* or *double*.

offset must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*, *local*, *private*) to *half*.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	175	*<id>* *data*	*<id>* *offset*	*<id>* *p*

vstore_half_r

Converts *data* float or double value to a half value and then write the converted value to the address (*p* + *offset*). The address computed as (*p* + *offset*) must be 16-bit aligned.

This function uses *mode* rounding mode when converting *data* to a half value.

data must be *float* or *double*.

offset must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*, *local*, *private*) to *half*.

9	44	*<id> Result Type*	*Result <id>*	extended instruc- tions set *<id>*	176	*<id> data*	*<id> offset*	*<id> p*	FP Rounding Mode *mode*

vstore_halfn

Converts *data* vector of float or vector of double values to a vector of half values and then write the converted value to the address (*p* + (*offset* * compCountOf(*data*))), where compCountOf(*data*) is equal to the component count of the vector *data*.

The address computed as (*p* + (*offset* * compCountOf(*data*))) must be 16-bit aligned.

This function uses the default rounding mode when converting *data* to a vector of half values. The default rounding mode is round to nearest even.

offset must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*, *local*, *private*) to *half*.

data must be *vector(2,3,4,8,16)* of *float* or *double* values.

8	44	*<id> Result Type*	*Result <id>*	extended instructions set *<id>*	177	*<id> data*	*<id> offset*	*<id> p*

vstore_halfn_r

Converts *data* vector of float or vector of double values to a vector of half values and then write the converted value to the address (*p* + (*offset* * compCountOf(*data*))), where compCountOf(*data*) is equal to the component count of the vector *data*.

The address computed as (*p* + (*offset* * compCountOf(*data*))) must be 16-bit aligned.

This function uses *mode* rounding mode when converting *data* to a half value.

offset must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*, *local*, *private*) to *half*.

data must be *vector(2,3,4,8,16)* of *float* or *double* values.

9	44	*<id>* *Result Type*	*Result* *<id>*	extended instructions set *<id>*	178	*<id>* *data*	*<id>* *offset*	*<id>* *p*	FP Rounding Mode *mode*

vloada_halfn

Reads a half vector value from the address (*p* + (*offset* * *n*)) and converts it to a float vector return value. The address computed as (*p* + (*offset* * *n*)) must be (2 * *n*) bytes aligned, when n = 2,4,8,16; For *n* = 3, the function returns a vector of 3 float values from the address (*p* + (*offset* * 4)). The address computed as (*p* + (*offset* * 4)) must be 8-bytes aligned

offset must be *size_t*.

p must be a *pointer*(*global*, *local*, *private*, *constant*) to *half*.

Result Type must be *vector(2,3,4,8,16)* of *float* values.

Result Type component count must be equal to *n*.

n must be 2,3,4,8 or 16.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	179	*<id>* *offset*	*<id>* *p*	Literal Number *n*

vstorea_halfn

Converts *data* vector of float or vector of double values to a vector of half values and then write the converted value to the address (*p* + (*offset* * compCountOf(*data*))), where compCountOf(*data*) is equal to the component count of the vector *data*.

The address computed as (*p* + (*offset* * compCountOf(*data*))) must be (2 * compCountOf(*data*)) bytes aligned, when n = 2,4,8,16; For *n* = 3, the function returns a vector of 3 float values from the address (*p* + (*offset* * 4)). The address computed as (*p* + (*offset* * 4)) must be 8-bytes aligned.

This function uses the default rounding mode when converting *data* to a vector of half values. The default rounding mode is round to nearest even.

offset must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*, *local*, *private*) to *half*.

data must be *vector(2,3,4,8,16)* of *float* or *double* values.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	180	*<id>* *data*	*<id>* *offset*	*<id>* *p*

vstorea_halfn_r

Converts *data* vector of float or vector of double values to a vector of half values and then write the converted value to the address (*p* + (*offset* * compCountOf(*data*))), where compCountOf(*data*) is equal to the component count of the vector *data*.

The address computed as (*p* + (*offset* * compCountOf(*data*))) must be (2 * compCountOf(*data*)) bytes aligned, when n = 2,4,8,16; For *n* = 3, the function returns a vector of 3 float values from the address (*p* + (*offset* * 4)). The address computed as (*p* + (*offset* * 4)) must be 8-bytes aligned.

This function uses *mode* rounding mode when converting *data* to a vector of half values.

offset must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*, *local*, *private*) to *half*.

data must be *vector(2,3,4,8,16)* of *float* or *double* values.

9	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	181	*<id>* *data*	*<id>* *offset*	*<id>* *p*	FP Rounding Mode *mode*

2.7   Miscellaneous Vector instructions

This section describes additional vector instructions.

shuffle

Construct a permutation of components from *x* vector value, returning a vector value with the same component type as *x* and component component count that is the same as *shuffle mask*.

In this function, only the *ilogb*(2 *m* -1) least significant bits of each mask element are considered, where *m* is equal to the component count of *x*.

shuffle mask operand specifies, for each component in the result vector, which component of *x* it gets.

The size of each component in *shuffle mask* must match the size of each component in *Result Type*.

Result Type must have the same component type as *x* and component count as *shuffle mask*.

shuffle mask must be *vector(2,4,8,16)* of *integer* values.

Result Type and *x* must be *vector(2,4,8,16)* of *floating-point* or *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	182	*<id>* *x*	*<id>* *shuffle mask*

shuffle2

Construct a permutation of components from *x* and *y* vector values, returning a vector value with the same component type as *x* and *y* and component count that is the same as *shuffle mask*.

In this function, only the *ilogb*(2 *m* - 1) + 1 least significant bits of each mask component are considered, where *m* is equal to the component count of *x* and *y*.

shuffle mask operand specifies, for each component in the result vector, which component of *x* or *y* it gets. Where component count begins with *x* and then proceeds to *y*.

x and *y* must be of the same type.

The size of each component in *shuffle mask* must match the size of each component in *Result Type*.

Result Type must have the same component type as *x* and component count as *shuffle mask*.

shuffle mask must be *vector(2,4,8,16)* of *integer* values.

Result Type,*x* and *y* must be *vector(2,4,8,16)* of *floating-point* or *integer* values.

All of the operands, including the *Result Type* operand, must be of the same type.

8	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	183	*<id>* *x*	*<id>* *y*	*<id>* *shuffle mask*

2.8 Misc instructions

This section describes additional miscellaneous instructions.

---

printf

The *printf* extended instruction writes output to an implementation-defined stream such as stdout under control of the string pointed to by format that specifies how subsequent arguments are converted for output. If there are insufficient arguments for the format, the behavior is undefined. If the format is exhausted while arguments remain, the excess arguments are evaluated (as always) but are otherwise ignored. The printf function returns when the end of the format string is encountered

printf returns 0 if it was executed successfully and -1 otherwise

Result Type must be *i32*.

format must be **OpString**.

6 + vari-able	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	184	*<id>* *format*	*<id>, <id>,* … *additional arguments*

---

prefetch

Prefetch *num_elements* * size in bytes of the type pointed by *p*, into the global cache. The prefetch instruction is applied to a work-item in a work-group and does not affect the functional behavior of the kernel.

num_elements must be *size_t*.

Result Type must be *void*.

p must be a *pointer*(*global*) to *floating-point*, *integer* or *vector(2,3,4,8,16)* of *floating-point*, *integer* values.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	185	*<id>* *num_elements*	*<id>* *p*

2.9 Image functions

The instructions defined in this section can only be used with image memory objects. An image memory object can be accessed by specific function calls that read from and/or write to specific locations in the image.

2.9.1 Image encoding

The following list denotes the different valid OpTypeSampler encodings of image objects.

---

image1d

A 1D image

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 0	*Image* *Type* 1	*Array* 0	*Depth* 0	*Sample* 0	Access Qualifier *qualifier*

image1dBuffer

A 1D image created from a buffer object.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 5	*Image* *Type* 1	*Array* 0	*Depth* 0	*Sample* 0	Access Qualifier *qualifier*

image1dArray

A 1D image array.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 0	*Image* *Type* 1	*Array* 1	*Depth* 0	*Sample* 0	Access Qualifier *qualifier*

image2d

A 2D image.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 0	*Depth* 0	*Sample* 0	Access Qualifier *qualifier*

image2dArray

A 2D image array.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 1	*Depth* 0	*Sample* 0	Access Qualifier *qualifier*

image2dDepth

A 2D depth image.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 0	*Depth* 1	*Sample* 0	Access Qualifier *qualifier*

image2dArrayDepth

A 2D depth image array.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 1	*Depth* 1	*Sample* 0	Access Qualifier *qualifier*

image2dMsaa

A 2D multi-sample color image.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 0	*Depth* 0	*Sample* 1	Access Qualifier *qualifier*

image2dArrayMsaa

A 2D multi-sample color image array.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 1	*Depth* 0	*Sample* 1	Access Qualifier *qualifier*

image2dMsaaDepth

A 2D multi-sample depth image.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 0	*Depth* 1	*Sample* 1	Access Qualifier *qualifier*

image2dArrayMsaaDepth

A 2D multi-sample depth image array.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 1	*Image* *Type* 1	*Array* 1	*Depth* 1	*Sample* 1	Access Qualifier *qualifier*

image3d

A 1D image created from a buffer object.

9	14	*Result* *<id>*	*Sampled* *Type <0>*	Dim 2	*Image* *Type* 1	*Array* 0	*Depth* 0	*Sample* 0	Access Qualifier *qualifier*

2.9.2 Sampler encoding

A SPIR-V *sampler* object is encoded via the **OpTypeSampler** instruction in the following way:

sampler

An image sampler object.

8	14	*Result <id>*	*Sampled* *Type <0>*	Dim 0	*Image Type* 2	*Array* 0	*Depth* 0	*Sample* 0

In addition, it is possible to define a constant *sampler* using the **OpConstantSampler**.

2.9.3   Image format encoding

Every image memory object has a format. An image format is a combination of *channel order* and *channel data type*. The *channel order* specifies the number of channels and the channel layout i.e.the memory layout in which channels are stored in the image. The *channel data type* describes the size of the channel data type.

ImageChannelOrder	
4272	**R**
4273	**A**
4274	**RG**
4275	**RA**
4276	**RGB**
4277	**RGBA**
4278	**BGRA**
4279	**ARGB**
4280	**INTENSITY**
4281	**LUMINANCE**
4282	**Rx**
4283	**RGx**
4284	**RGBx**
4285	**DEPTH**
4286	**DEPTH STENCIL**
4287	**sRGB**
4288	**sRGBx**
4289	**sRGBA**
4290	**sBGRA**

ImageChannelType	
4304	**SNORM INT8**
4305	**SNORM INT16**
4306	**UNORM INT8**
4307	**UNORM_INT16**
4308	**UNORM SHORT 565**
4309	**UNORM SHORT 555**
4310	**UNORM INT 101010**
4311	**SIGNED INT8**
4312	**SIGNED INT16**
4313	**SIGNED INT32**
4314	**UNSIGNED INT8**
4315	**UNSIGNED INT16**
4316	**UNSIGNED INT32**
4317	**HALF FLOAT**
4318	**FLOAT**
4319	**UNORM INT24**

2.9.4   Image read functions

This section describes the list of instructions that allow reading from image memory objects.

---

read_imagef

Use the coordinate specified by *coords* and the *sampler* object specified by *s* to do an element lookup to the image object specified by *img*.

This function returns floating-point values in the range [0.0 . . . 1.0] for *image* objects created with *channel data type* set to one of the pre-defined packed formats or **UNORM INT8**, or **UNORM INT16**.

This function returns floating-point values in the range[-1.0 . . . 1.0] for *image* objects created with *channel data type* set to **SNORM INT8**, or **SNORM INT16**.

This function returns floating-point values for *image* objects created with *channel data type* set to **HALF FLOAT**, or **FLOAT**.

When called with i32 coordinates the sampler object must be defined with a filter mode set to **Nearest**, coordinates set to non-parametric coordinates and addressing mode set to **ClampToEdge**, **Clamp** or **None**; otherwise the values returned are undefined.

Values returned by this function for image objects with *channel data type* which is not specified in the description above are undefined.

Result Type must be *float* or *vector(4)* of *float* values.

coords must be *float* or *i32* or *vector(2,4)* of *float* or *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth*, *image2dMsaa*, *image2dArrayMsaa*, *image2dMsaaDepth*, *image2dArrayMsaaDepth* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

When *img* is a *image1d*, *coords* must be *float* or *i32*.

When *img* is a *image2d*, *image2dDepth*, *image2dMsaa* or *image2dMsaaDepth*, *coords* must be *vector(2)* of *float* or *i32* values.

When *img* is a *image1dArray*, *coords* must be *vector(2)* of *i32* values. The second component of *coords* is used to identify the image in the array

When *img* is a *image2dArray*, *image2dArrayDepth*, *image2dArrayMsaa* or *image2dArrayMsaaDepth*, *coords* must be *vector(4)* of *i32* values. The third component of *coords* is used to identify the image in the array, while the fourth component is ignored.

When *img* is a *image3d*, *coords* must be *vector(4)* of *float* or *i32* values. The fourth component of *coords* is ignored.

Result Type must be a *float* when *img* is a *image2dArrayDepth*, *image2dDepth*, *image2dMsaaDepth* or *image2dArrayMsaaDepth*, and *vector(4)* of *float* values when *img* is on of the remaining valid image types for this instruction.

| 8 | 44 | *<id>* *Result Type* | Result *<id>* | extended instructions set *<id>* | 111 | *<id>* *img* | *<id>* *s* | *<id>* *coords* |

read_imagei

Use the coordinate specified by *coords* and the *sampler* object specified by *s* to do an element lookup to the image object specified by *img*.

This function returns a non-parametric *i32* integer value.

This function can only be used if *img* image object *channel data type* is set to **SIGNED INT8**, **SIGNED INT16** or **SIGNED INT32**. If the *channel data type* is not one of these values, the values returned by read_imagei are undefined.

The sampler object must be defined with a filter mode set to **Nearest**, coordinates set to non-parametric coordinates and addressing mode set to **ClampToEdge**, **Clamp** or **None**; otherwise the values returned are undefined.

Result Type must be *vector(4)* of *i32* values.

coords must be *float* or *i32* or *vector(2,4)* of *float* or *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray*, *image2dMsaa*, *image2dArrayMsaa* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

When *img* is a *image1d*, *coords* must be *float* or *i32*.

When *img* is a *image2d*, *image2dDepth*, *image2dMsaa* or *image2dMsaaDepth*, *coords* must be *vector(2)* of *float* or *i32* values.

When *img* is a *image1dArray*, *coords* must be *vector(2)* of *i32* values. The second component of *coords* is used to identify the image in the array

When *img* is a *image2dArray*, *image2dArrayDepth*, *image2dArrayMsaa* or *image2dArrayMsaaDepth*, *coords* must be *vector(4)* of *i32* values. The third component of *coords* is used to identify the image in the array, while the fourth component is ignored.

When *img* is a *image3d*, *coords* must be *vector(4)* of *float* or *i32* values. The fourth component of *coords* is ignored.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	112	*<id>* *img*	*<id>* *s*	*<id>* *coords*

read_imageui

Use the coordinate specified by *coords* and the *sampler* object specified by *s* to do an element lookup to the image object specified by *img*.

This function returns a non-parametric *i32* integer value.

This function can only be used if *img* image object *channel data type* is set to **UNSIGNED INT8**, **UNSIGNED INT16** or **UNSIGNED INT32**. If the *channel data type* is not one of these values, the values returned by read_imagei are undefined.

The sampler object must be defined with a filter mode set to **Nearest**, coordinates set to non-parametric coordinates and addressing mode set to **ClampToEdge**, **Clamp** or **None**; otherwise the values returned are undefined.

Result Type must be *vector(4)* of *i32* values.

coords must be *float* or *i32* or *vector(2,4)* of *float* or *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray*, *image2dMsaa*, *image2dArrayMsaa* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

When *img* is a *image1d*, *coords* must be *float* or *i32*.

When *img* is a *image2d*, *image2dDepth*, *image2dMsaa* or *image2dMsaaDepth*, *coords* must be *vector(2)* of *float* or *i32* values.

When *img* is a *image1dArray*, *coords* must be *vector(2)* of *i32* values. The second component of *coords* is used to identify the image in the array

When *img* is a *image2dArray*, *image2dArrayDepth*, *image2dArrayMsaa* or *image2dArrayMsaaDepth*, *coords* must be *vector(4)* of *i32* values. The third component of *coords* is used to identify the image in the array, while the fourth component is ignored.

When *img* is a *image3d*, *coords* must be *vector(4)* of *float* or *i32* values. The fourth component of *coords* is ignored.

| 8 | 44 | *<id>* *Result Type* | *Result <id>* | extended instructions set *<id>* | 113 | *<id>* *img* | *<id>* *s* | *<id>* *coords* |

read_imageh

Use the coordinate specified by *coords* and the *sampler* object specified by *s* to do an element lookup to the image object specified by *img*.

This function returns half precision floating-point values in the range [0.0 ... 1.0] for *image* objects created with *channel data type* set to one of the pre-defined packed formats or **UNORM INT8**, or **UNORM INT16**.

This function returns half precision floating-point values in the range[-1.0 ... 1.0] for *image* objects created with *channel data type* set to **SNORM INT8**, or **SNORM INT16**.

This function returns half precision floating-point values for *image* objects created with *channel data type* set to **HALF FLOAT**, or **FLOAT**.

When called with i32 coordinates the sampler object must be defined with a filter mode set to **Nearest**, coordinates set to non-parametric coordinates and addressing mode set to **ClampToEdge**, **Clamp** or **None**; otherwise the values returned are undefined.

Values returned by this function for image objects with *channel data type* which is not specified in the description above are undefined.

Result Type must be *half* or *vector(4)* of *half* values.

coords must be *float* or *i32* or *vector(2,4)* of *float* or *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

When *img* is a *image1d*, *coords* must be *float* or *i32*.

When *img* is a *image2d*, *coords* must be *vector(2)* of *float* or *i32* values.

When *img* is a *image1dArray*, *coords* must be *vector(2)* of *i32* values. The second component of *coords* is used to identify the image in the array

When *img* is a *image2dArray*, *coords* must be *vector(4)* of *i32* values. The third component of *coords* is used to identify the image in the array, while the fourth component is ignored.

When *img* is a *image3d*, *coords* must be *vector(4)* of *float* or *i32* values. The fourth component of *coords* is ignored.

Result Type must be a *half* when *img* is a *image2dArrayDepth*, *image2dDepth*, *image2dMsaaDepth* or *image2dArrayMsaaDepth*, and *vector(4)* of *half* values when *img* is on of the remaining valid image types for this instruction.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	114	*<id>* *img*	*<id>* *s*	*<id>* *coords*

read_imagef_samplerless

Use the coordinate specified by *coords* to do an element lookup to the image object specified by *img*. This function behaves exactly as the corresponding read_imagef function that take integer coordinates and a sampler with filter mode set to **Nearest**, non-parametric coordinates and addressing mode set to **None**.

Result Type must be *float* or *vector(4)* of *float* values.

coords must be *i32* or *vector(2,4)* of *i32* values.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	115	*<id>* *img*	*<id>* *coords*

read_imagei_samplerless

Use the coordinate specified by *coords* to do an element lookup to the image object specified by *img*. This function behaves exactly as the corresponding read_imagei function that take integer coordinates and a sampler with filter mode set to **Nearest**, non-parametric coordinates and addressing mode set to **None**.

Result Type must be *vector(4)* of *i32* values.

coords must be *i32* or *vector(2,4)* of *i32* values.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray*, *image2dMsaa*, *image2dArrayMsaa* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	116	*<id>* *img*	*<id>* *coords*

read_imageui_samplerless

Use the coordinate specified by *coords* to do an element lookup to the image object specified by *img*. This function behaves exactly as the corresponding read_imageui function that take integer coordinates and a sampler with filter mode set to **Nearest**, non-parametric coordinates and addressing mode set to **None**.

Result Type must be *vector(4)* of *i32* values.

coords must be *i32* or *vector(2,4)* of *i32* values.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

7	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	117	*<id>* *img*	*<id>* *coords*

read_imageh_samplerless

Use the coordinate specified by *coords* to do an element lookup to the image object specified by *img*. This function behaves exactly as the corresponding read_imageh function that take integer coordinates and a sampler with filter mode set to **Nearest**, non-parametric coordinates and addressing mode set to **None**.

Result Type must be *vector(4)* of *half* values.

coords must be *i32* or *vector(2,4)* of *i32* values.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

7	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	118	*<id>* *img*	*<id>* *coords*

read_imagef_mipmap_lod

Use the coordinate specified by *coords*, and the sampler object specified by *s* to do an element lookup in the mip-level specified by *lod* in the image object specified by *img*.

Result Type must be *float* or *vector(4)* of *float* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

s must be set to use parametric coordinates.

lod is clamped to the minimum of (actual number of mip-levels - 1) in the image or value specified for CL_SAMPLER_LOD_MAX.

When *img* type is *image2d*:

- *coords* must be a *vector(2)* of *float* values.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image2dArray*:

- *coords* must be a *vector(4)* of *float* values.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image1d*:

- *coords* must be a *float*.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image1dArray*:

- *coords* must be a *vector(2)* of *float* values.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image3d*:

- *coords* must be a *vector(4)* of *float* values.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image2dDepth*:

- *coords* must be a *vector(2)* of *float* values.

- *Result Type* must be a *float*.

When *img* type is *image2dArrayDepth*:

- *coords* must be a *vector(4)* of *float* values.

- *Result Type* must be a *float*.

9	44	*<id>* *Result* *Type*	*Result* *<id>*	extended instruc- tions set *<id>*	123	*<id>* *img*	*<id>* *s*	*<id>* *coords*	*<id>* *lod*

---

read_imagei_mipmap_lod

Use the coordinate specified by *coords*, and the sampler object specified by *s* to do an element lookup in the mip-level specified by *lod* in the image object specified by *img*.

Result Type must be *vector(4)* of *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

s must be set to use parametric coordinates.

lod is clamped to the minimum of (actual number of mip-levels - 1) in the image or value specified for CL_SAMPLER_LOD_MAX.

When *img* type is *image2d*:

- *coords* must be a *vector(2)* of *float* values.

- *lod* must be a *float*.

When *img* type is *image2dArray*:

- *coords* must be a *vector(4)* of *float* values.

When *img* type is *image1d*:

- *coords* must be a *float*.

When *img* type is *image1dArray*:

- *coords* must be a *vector(2)* of *float* values.

When *img* type is *image3d*:

- *coords* must be a *vector(4)* of *float* values.

9	44	*<id>* *Result* *Type*	*Result* *<id>*	extended instruc- tions set *<id>*	124	*<id>* *img*	*<id>* *s*	*<id>* *coords*	*<id>* *lod*

read_imageui_mipmap_lod

Use the coordinate specified by *coords*, and the sampler object specified by *s* to do an element lookup in the mip-level specified by *lod* in the image object specified by *img*.

Result Type must be *vector(4)* of *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

s must be set to use parametric coordinates.

lod is clamped to the minimum of (actual number of mip-levels - 1) in the image or value specified for CL_SAMPLER_LOD_MAX.

When *img* type is *image2d*:

- *coords* must be a *vector(2)* of *float* values.

When *img* type is *image2dArray*:

- *coords* must be a *vector(4)* of *float* values.

When *img* type is *image1d*:

- *coords* must be a *float*.

When *img* type is *image1dArray*:

- *coords* must be a *vector(2)* of *float* values.

When *img* type is *image3d*:

- *coords* must be a *vector(4)* of *float* values.

9	44	*<id>* *Result* *Type*	*Result* *<id>*	extended instruc- tions set *<id>*	125	*<id>* *img*	*<id>* *s*	*<id>* *coords*	*<id>* *lod*

read_imagef_mipmap_gradient

Use the gradients *grad_x* and *grad_y*, the coordinates specified by *coords*, and the sampler object specified by *s* to do an element lookup in the computed mip-level in the image object specified by *img*.

Result Type must be *float* or *vector(4)* of *float* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

s must be set to use parametric coordinates.

When *img* type is *image2d*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image2dArray*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image1d*:

- *coords* must be a *float*.

- *grad_x* and *grad_y* must be a *float*.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image1dArray*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *float*.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image3d*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(4)* of *float* values.

- *Result Type* must be a *vector(4)* of *float* values.

When *img* type is *image2dDepth*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

- *Result Type* must be a *float*.

When *img* type is *image2dArrayDepth*:

10	44	*<id>* *Result* *Type*	*Result* *<id>*	extended instruc- tions set *<id>*	126	*<id>* *img*	*<id>* *s*	*<id>* *coords*	*<id>* *grad_x*	*<id>* *grad_y*

read_imagei_mipmap_gradient

Use the gradients *grad_x* and *grad_y*, the coordinates specified by *coords*, and the sampler object specified by *s* to do an element lookup in the computed mip-level in the image object specified by *img*.

Result Type must be *vector(4)* of *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

s must be set to use parametric coordinates.

When *img* type is *image2d*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

When *img* type is *image2dArray*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

When *img* type is *image1d*:

- *coords* must be a *float*.

- *grad_x* and *grad_y* must be a *float*.

When *img* type is *image1dArray*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *float*.

When *img* type is *image3d*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(4)* of *float* values.

When *img* type is *image2dDepth*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

When *img* type is *image2dArrayDepth*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

10	44	*<id> Result Type*	*Result <id>*	extended instruc- tions set *<id>*	127	*<id> img*	*<id> s*	*<id> coords*	*<id> grad_x*	*<id> grad_y*

10	44	*<id> Result Type*	*Result <id>*	extended instruc- tions set *<id>*	127	*<id> img*	*<id> s*	*<id> coords*	*<id> grad_x*	*<id> grad_y*

read_imageui_mipmap_gradient

Use the gradients *grad_x* and *grad_y*, the coordinates specified by *coords*, and the sampler object specified by *s* to do an element lookup in the computed mip-level in the image object specified by *img*.

Result Type must be *vector(4)* of *i32* values.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with ReadOnly or ReadWrite access qualifier.

s must be *sampler* value.

s must be set to use parametric coordinates.

When *img* type is *image2d*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

When *img* type is *image2dArray*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

When *img* type is *image1d*:

- *coords* must be a *float*.

- *grad_x* and *grad_y* must be a *float*.

When *img* type is *image1dArray*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *float*.

When *img* type is *image3d*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(4)* of *float* values.

When *img* type is *image2dDepth*:

- *coords* must be a *vector(2)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

When *img* type is *image2dArrayDepth*:

- *coords* must be a *vector(4)* of *float* values.

- *grad_x* and *grad_y* must be a *vector(2)* of *float* values.

10	44	<id> Result Type	Result <id>	extended instruc- tions set <id>	128	<id> img	<id> s	<id> coords	<id> grad_x	<id> grad_y

2.9.5 Image write functions

This section describes the list of instructions that allow writing to image memory objects.

10	44	<id> Result Type	Result <id>	extended instruc- tions set <id>	128	<id> img	<id> s	<id> coords	<id> grad_x	<id> grad_y

write_imagef

Write *value* to the coordinates specified by *coords* to the image object specified by *img*. The write happens only after the data in *value* is converted to the appropraite *img* image *channel data type*. *coords* are considered to be non-parametric coordinates.

Result Type must be *void*.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with WriteOnly or ReadWrite access qualifier.

When *img* is a *image2d*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**, **FLOAT**.

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image height - 1) respectively.

- *value* is a *vector(4)* of *float* values.

When *img* is a *image2dArray*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**, **FLOAT**.

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width - 1), (0 . . . image height - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

- *value* is a *vector(4)* of *float* values.

When *img* is a *image1d* or *image1dBuffer*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**, **FLOAT**.

- *coords* is a *i32*, and is in the range (0 . . . image width - 1)

- *value* is a *vector(4)* of *float* values.

When *img* is a *image1dArray*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**, **FLOAT**.

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image number of layers - 1) respectively.

- *value* is a *vector(4)* of *float* values

When *img* is a *image2dDepth*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM INT16**, **UNORM INT24**, **FLOAT**.

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image height - 1) respectively.

- *value* is a *float*.

When *img* is a *image2dArrayDepth*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM INT16**, **UNORM INT24**, **FLOAT**.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	119	*<id>* *img*	*<id>* *coords*	*<id>* *value*

---

write_imagei

Write *value* to the coordinates specified by *coords* to the image object specified by *img*. The write happens only after the data in *value* is converted to the appropraite *img* image *channel data type*. *value* component type is considered to be a signed integer. *coords* are considered to be non-parametric coordinates.

Result Type must be *void*.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with WriteOnly or ReadWrite access qualifier.

The *channel data type* of *img* must be set to **SIGNED INT8**, **SIGNED INT16**, **SIGNED INT32**.

When *img* is a *image2d*:

- *coords* must be a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image height - 1) respectively.

- *value* must be a *vector(4)* of *i32* values.

When *img* is a *image2dArray*:

- *coords* must be a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width - 1), (0 . . . image height - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

- *value* must be a *vector(4)* of *i32* values.

When *img* is a *image1d* or *image1dBuffer*:

- *coords* must be a *i32*, and is in the range (0 . . . image width - 1)

- *value* must be a *vector(4)* of *i32* values.

When *img* is a *image1dArray*:

- *coords* must be a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image number of layers - 1) respectively.

- *value* must be a *vector(4)* of *i32* values

When *img* is a *image3d*:

- *coords* must be a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width - 1), (0 . . . image height - 1), (0 . . . image depth - 1) respectively. The fourth component is ignored.

- *value* must be a *vector(4)* of *i32* values.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	120	*<id>* *img*	*<id>* *coords*	*<id>* *value*

write_imageui

Write *value* to the coordinates specified by *coords* to the image object specified by *img*. The write happens only after the data in *value* is converted to the appropraite *img* image *channel data type*. *value* component type is considered to be an unsigned integer. *coords* are considered to be non-parametric coordinates.

Result Type must be *void*.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with WriteOnly or ReadWrite access qualifier.

The *channel data type* of *img* must be set to **UNSIGNED INT8**, **UNSIGNED INT16**, **UNSIGNED INT32**.

When *img* is a *image2d*:

- *coords* must be a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image height - 1) respectively.

- *value* must be a *vector(4)* of *i32* values.

When *img* is a *image2dArray*:

- *coords* must be a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width - 1), (0 . . . image height - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

- *value* must be a *vector(4)* of *i32* values.

When *img* is a *image1d* or *image1dBuffer*:

- *coords* must be a *i32*, and is in the range (0 . . . image width - 1)

- *value* must be a *vector(4)* of *i32* values.

When *img* is a *image1dArray*:

- *coords* must be a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image number of layers - 1) respectively.

- *value* must be a *vector(4)* of *i32* values

When *img* is a *image3d*:

- *coords* must be a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width - 1), (0 . . . image height - 1), (0 . . . image depth - 1) respectively. The fourth component is ignored.

- *value* must be a *vector(4)* of *i32* values.

8	44	*<id>* *Result Type*	Result *<id>*	extended instructions set *<id>*	121	*<id>* *img*	*<id>* *coords*	*<id>* *value*

write_imageh

Write *value* to the coordinates specified by *coords* to the image object specified by *img*. The write happens only after the data in *value* is converted to the appropraite *img* image *channel data type*. *coords* are considered to be non-parametric coordinates.

Result Type must be *void*.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with WriteOnly or ReadWrite access qualifier.

When *img* is a *image2d*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**.

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image height - 1) respectively.

- *value* is a *vector(4)* of *half* values.

When *img* is a *image2dArray*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**.

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width - 1), (0 . . . image height - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

- *value* is a *vector(4)* of *half* values.

When *img* is a *image1d* or *image1dBuffer*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**.

- *coords* is a *i32*, and is in the range (0 . . . image width - 1)

- *value* is a *vector(4)* of *half* values.

When *img* is a *image1dArray*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**.

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width - 1), (0 . . . image number of layers - 1) respectively.

- *value* is a *vector(4)* of *half* values

When *img* is a *image3d*, the behavior of the function is undefined unless:

- The *channel data type* of *img* is set to **UNORM SHORT 565**,**UNORM SHORT 555**, **UNORM INT 101010**, **UNORM INT8**, **SNORM INT8**, **UNORM INT16**, **SNORM INT16**, **HALF FLOAT**.

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width - 1), (0 . . . image height - 1), (0 . . . image depth - 1) respectively. The fourth component is ignored.

- *value* is a *vector(4)* of *half* values.

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	122	*<id>* *img*	*<id>* *coords*	*<id>* *value*

8	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	122	*<id>* *img*	*<id>* *coords*	*<id>* *value*

---

write_imagef_mipmap_lod

Write *value* to the coordinates specified by *coords* in the mip-level specified by *lod* to the image object specified by *img*. The write happens only after the data in *value* is converted to the appropraite *img* image *channel data type*. *coords* are considered to be non-parametric coordinates.

Result Type must be *void*.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with WriteOnly or ReadWrite access qualifier.

The behavior of the function is undefined unless *lod* value is in the range (0 . . . number of mip-levels in the image - 1).

When *img* is a *image2d*, the behavior of the function is undefined unless:

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1) respectively.

- *value* is a *vector(4)* of *float* values.

When *img* is a *image2dArray*, the behavior of the function is undefined unless:

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

- *value* is a *vector(4)* of *float* values.

When *img* is a *image1d* or *image1dBuffer*, the behavior of the function is undefined unless:

- *coords* is a *i32*, and is in the range (0 . . . image width of the mip-level specified by *lod* - 1)

- *value* is a *vector(4)* of *float* values.

When *img* is a *image1dArray*, the behavior of the function is undefined unless:

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image number of layers - 1) respectively.

- *value* is a *vector(4)* of *float* values.

When *img* is a *image2dDepth*, the behavior of the function is undefined unless:

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width of the mip-level specified by *lod*- 1), (0 . . . image height of the mip-level specified by *lod*- 1) respectively.

- *value* is a *float*.

When *img* is a *image2dArrayDepth*, the behavior of the function is undefined unless:

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip - level specified by *lod* - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

- *value* is a *float*.

When *img* is a *image3d*, the behavior of the function is undefined unless:

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1), (0 . . . image depth of the mip-level specified by *lod* - 1) respectively. The fourth component is ignored.

- *value* is a *vector(4)* of *float* values.

9	44	*<id>* *Result* *Type*	*Result* *<id>*	extended instruc- tions set *<id>*	129	*<id>* *img*	*<id>* *coords*	*<id>* *lod*	*<id>* *value*

---

write_imagei_mipmap_lod

Write *value* to the coordinates specified by *coords* in the mip-level specified by *lod* to the image object specified by *img*. The write happens only after the data in *value* is converted to the appropraite *img* image *channel data type*. *coords* are considered to be non-parametric coordinates. *value* component type is treated as signed integer.

Result Type must be *void*.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with WriteOnly or ReadWrite access qualifier.

The behavior of the function is undefined unless *lod* value is in the range (0 . . . number of mip-levels in the image - 1).

When *img* is a *image2d*, the behavior of the function is undefined unless:

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1) respectively.

When *img* is a *image2dArray*, the behavior of the function is undefined unless:

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

When *img* is a *image1d* or *image1dBuffer*, the behavior of the function is undefined unless:

- *coords* is a *i32*, and is in the range (0 . . . image width of the mip-level specified by *lod* - 1)

When *img* is a *image1dArray*, the behavior of the function is undefined unless:

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image number of layers - 1) respectively.

When *img* is a *image3d*, the behavior of the function is undefined unless:

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1), (0 . . . image depth of the mip-level specified by *lod* - 1) respectively. The fourth component is ignored.

9	44	*<id>* *Result* *Type*	*Result* *<id>*	extended instruc- tions set *<id>*	130	*<id>* *img*	*<id>* *coords*	*<id>* *lod*	*<id>* *value*

write_imageui_mipmap_lod

Write *value* to the coordinates specified by *coords* in the mip-level specified by *lod* to the image object specified by *img*. The write happens only after the data in *value* is converted to the appropraite *img* image *channel data type*. *coords* are considered to be non-parametric coordinates. *value* component type is treated as unsigned integer.

Result Type must be *void*.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray* or *image3d* value, with WriteOnly or ReadWrite access qualifier.

The behavior of the function is undefined unless *lod* value is in the range (0 . . . number of mip-levels in the image - 1).

When *img* is a *image2d*, the behavior of the function is undefined unless:

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1) respectively.

When *img* is a *image2dArray*, the behavior of the function is undefined unless:

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1), (0 . . . image number of layers - 1) respectively. The fourth component is ignored.

When *img* is a *image1d* or *image1dBuffer*, the behavior of the function is undefined unless:

- *coords* is a *i32*, and is in the range (0 . . . image width of the mip-level specified by *lod* - 1)

When *img* is a *image1dArray*, the behavior of the function is undefined unless:

- *coords* is a *vector(2)* of *i32* values, where the first and second components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image number of layers - 1) respectively.

When *img* is a *image3d*, the behavior of the function is undefined unless:

- *coords* is a *vector(4)* of *i32* values, where the first, second and third components are in the range (0 . . . image width of the mip-level specified by *lod* - 1), (0 . . . image height of the mip-level specified by *lod* - 1), (0 . . . image depth of the mip-level specified by *lod* - 1) respectively. The fourth component is ignored.

9	44	*<id>* *Result Type*	*Result* *<id>*	extended instruc- tions set *<id>*	131	*<id>* *img*	*<id>* *coords*	*<id>* *lod*	*<id>* *value*

2.9.6 Image query functions

This section describes the list of instructions that provide information of image memory objects.

get_image_width

Return the width in pixels of the image object specified by *img*.

Result Type must be *i32*.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	132	*<id>* *img*

get_image_height

Return the height in pixels of the image object specified by *img*.

Result Type must be *i32*.

img must be *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	133	*<id>* *img*

get_image_depth

Return the depth in pixels of the image object specified by *img*.

Result Type must be *i32*.

img must be *image3d* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	134	*<id>* *img*

get_image_channel_data_type

Return the *channel data type* of the image object specified by *img*.

Result Type must be *i32*.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

Result Type must contain a value from ImageChannelType enumeration.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	135	*<id>* *img*

get_image_channel_order

Return the *channel order* of the image object specified by *img*.

Result Type must be *i32*.

img must be *image1d*, *image1dBuffer*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

Result Type must contain a value from ImageChannelOrder enumeration.

6	44	*<id>* Result Type	Result *<id>*	extended instructions set *<id>*	136	*<id>* img

get_image_dim

Return the dimensions of the image object specified by *img*.

Result Type must be *i32* or *vector(2,4)* of *i32* values.

img must be *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

Result Type' must be 'vector(2) of *i32* values when *img* is a *image2d*, *image2dArray*, *image2dArrayDepth* or *image2dDepth*. The width and height of the image are contained in the first and second components of the return value repectively.

Result Type' must be 'vector(4) of *i32* values when *img* is a *image3d*. The width, height and depth of the image are contained in the first, second and third components of the return value repectively. The fourth component is 0.

6	44	*<id>* Result Type	Result *<id>*	extended instructions set *<id>*	137	*<id>* img

get_image_array_size

Return the number of samples in the MSAA image object specified by *img*.

Result Type must be *i32*.

Result Type must be *size_t*.

img must be *image1dArray*, *image2dArray* or *image2dArrayDepth* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

img must be *image2dMsaa*, *image2dArrayMsaa*, *image2dMsaaDepth* or *image2dArrayMsaaDepth* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

7	44	*<id>* Result Type	Result *<id>*	extended instructions set *<id>*	138	*<id>* img	*<id>* img

get_image_num_mip_levels

Return the number of mip-levels of the image object specified by *img*.

Result Type must be *i32*.

img must be *image1d*, *image1dArray*, *image2d*, *image2dArray*, *image2dArrayDepth*, *image2dDepth* or *image3d* value, with ReadOnly, WriteOnly or ReadWrite access qualifier.

6	44	*<id>* *Result Type*	*Result <id>*	extended instructions set *<id>*	140	*<id>* *img*